



Document Visual Question Answering Using RAG

Submitted in partial fulfilment of the requirements for the degree of Bachelor of Science in Computing & Artificial Intelligence at the Artificial Intelligence Department, the Faculty of Computing & Artificial Intelligence, Helwan University

Supervised by:

Dr. Amr S. Ghoneim

June 2025

Abstract

Extracting accurate and relevant information from complex documents—such as medical records, financial reports, and legal contracts—poses a significant challenge in the field of Document Artificial Intelligence (Document AI). These documents are typically multi-page, visually dense, and rich in structured elements such as tables, charts, headers, and footnotes. As organizations increasingly depend on automated systems to interact with such data, the ability to ask natural language questions and receive accurate, context-aware answers—referred to as Document Visual Question Answering (DocVQA)—has become a critical requirement.

Unlike standard Visual Question Answering tasks that operate on natural images, DocVQA introduces unique challenges. Document understanding demands not only textual comprehension but also visual reasoning, as meaning is often derived from layout, position, and formatting. Moreover, enterprise documents frequently span several pages, creating long-context scenarios that exceed the input limitations of even the most advanced language or vision-language models. Traditional approaches struggle with this balance: systems that rely on Optical Character Recognition (OCR) and text-based retrieval often ignore visual context, leading to incorrect or incomplete answers. On the other hand, vision-language models that interpret document layout and visual structure are typically inefficient when dealing with long-form, multi-page inputs.

To address this dichotomy, this project introduces a hybrid framework for Document Visual Question Answering that combines the scalability of retrieval-based approaches with the visual awareness of multimodal understanding. The proposed solution retrieves the most relevant document segments by jointly considering textual and layout-based signals, and then passes this context to a vision-language generation module capable of producing fluent, accurate answers grounded in the document's structure and content. This architecture enables the system to handle documents of arbitrary length while maintaining an understanding of complex visual and structural cues.

By decoupling retrieval and generation, and leveraging both textual and visual modalities, the proposed framework effectively mitigates the limitations of previous systems. This makes it suitable for real-world applications where both the scale and structure of documents matter—such as medical question answering, contract analysis, and enterprise document automation. The system is evaluated on diverse document types and question categories, demonstrating improved performance in answering questions that require both cross-page context synthesis and visual reasoning. This work represents a step forward toward building AI assistants that can truly read, understand, and interact with documents as humans do.

Keywords

Document AI, Document Visual Question Answering (DocVQA), Long-Context Documents, Multimodal Models, Retrieval-Augmented Generation (RAG), Vision-Language Models (VLMs), M3socrag.

Acknowledgement

We would like to express our deepest gratitude to **Dr. Amr S. Ghoneim** for his exceptional mentorship and unwavering support throughout our graduation project, "**Document Visual Question Answering Using RAG**," at Helwan University. His profound knowledge, insightful guidance, and continuous encouragement were instrumental in steering our research toward success.

Dr. Ghoneim's expertise in **AI** played a crucial role in helping us shape the research direction and overcome technical challenges. His valuable feedback and commitment to excellence significantly enriched the quality of our work and deepened our understanding of the field.

We are also profoundly thankful to our families—especially our parents—for their endless emotional and financial support throughout our academic journey. Their encouragement, sacrifices, and unwavering belief in us have been the cornerstone of our growth and achievements.

Finally, we extend our sincere appreciation to the **Faculty of Computers and Artificial Intelligence** at Helwan University for fostering a rich academic environment. Their commitment to excellence has provided us with the tools and inspiration to pursue innovation and excellence in the field of computer science.

Table of Contents

Abstract.....	2
Keywords.....	3
Acknowledgement	4
List of Figures	7
List of Tables	7
List of Abbreviations	8
List of Equations.....	9
Glossary.....	10
Chapter 1 : An Introduction	12
1.1 Overview	14
1.2 Problem Statement.....	15
1.3 Scope and Objectives	26
Project Objectives	26
1.4 Report Organization (Structure)	27
1.5 Work Methodology	28
1. Experimental Research Framework	28
2. Agile-Inspired Development Process	29
3. SMART Goals Framework.....	30
1.6 Work Plan (Gantt chart).....	31
Chapter 2: Related Work (Literature Review).....	32
2.1 Background	32
Summary Points:	35
2.2 Literature Survey.....	35
2.3Analysis of Related Work	39
2.4 Summary and Research Gap	44
Chapter 3: The Proposed Solution	46
3.1 Solution Methodology	46
3.1.1Problem Scope Refinement	46
3.1.2 System Architecture Overview (M3DocRAG Framework)	47
3.1.3 Problem definition	48
3.1.4 Document Embedding	49
3.15 Page Retrieval	50
3.1.6 Question Answering.....	50
3.1.7 Models and Tools Selection	51
Embedding Methods Comparison: LAION CLIP, ColPaLi, and ColQwen	53

3.2.1 Functional Requirements	73
3.2.1 Non-Functional Requirements.....	73
4. Implementation and Experimental Setup.....	74
 4.1 Dataset Unification.....	74
 4.1.2 Unified Dataset Structure.....	75
 4.1.3 Dataset Unification Process.....	76
 4.1.4 Implementation Details.....	76
 4.2 Decoupling Retrieval and Generation in the RAG Pipeline.....	77
 4.2.1 Motivation and Architectural Overview	77
 4.2.2 Retrieval Task: Document Page Selection	78
 4.2.4 Design Patterns and Component Modularity.....	79
 4.2.5 Interoperability and Execution Flow	80
 4.2.6 End-to-End RAG Flow	80
5. Experimental Setup.....	80
 5.4 End-to-End Evaluation.....	83
 5.5 Generation Results (1k-shot Pilot).....	83
 Observations.....	84
 5.6 Retrieval Results (1k-shot Pilot)	85
 5.6.1 ColQwen (Enc-1B)	85
 5.6.2 ColPaLI (PaLI-5B-Enc).....	85
6. Conclusion & Future Work	86
 Key Lessons Learned	86
 Limitations.....	86
 Future Work.....	87
References (or Bibliography)	88

List of Figures

Figure 1 Example of question and generated answer for abstract images(VQA)	12
Figure 2 Example of question and generated answer for abstract images(DocVQA)	13
Figure 3Comparison of multi-modal document understanding pipelines. Previous works focus on (a) Single-page DocVQA that cannot.....	15
Figure 4 Large Language models architecture	16
Figure 5 Vision Language models architecture	18
Figure 6 Our reparameterization. We only train A and B	20
Figure 7 RAG architecture	21
Figure 8 RAG VS Fine Tuning	23
Figure 9 Information Retrieval.....	24
Figure 10 Neural re-ranking architectures.....	25
Figure 11 SMART Business Goals	31
Figure 12 Gantt chart	32
Figure 13 TextRAG (left) vs. VisRAG (right).....	51
Figure 14 ColPali simplifies document retrieval	56
Figure 15 (Left: Token Pooling) Relative performance degradation when reducing the number of stored embeddings per document. (Right: Interpretability) For each term in a user query.....	64

List of Tables

Table 1 Trade-Off: LLMs vs. VLMs in DocVQA.....	18
Table 2 Limitations of Traditional RAG for Document Visual Question Answering.....	22
Table 3 Core Trade-Offs in Current DocVQA Approaches	23
Table 4 Agile Process.....	30
Table 5 SMART Business Goals	31
Table 6 Work Plan	32
Table 7 Key Techniques in Document QA.....	34
Table 8 papers comparison	39
Table 9 papers advantges	44
Table 10 retrieval comparison	53
Table 11 These results of the evaluation	57
Table 12 Comprehensive evaluation of baseline models	58
Table 13 Comparison between ColPaLi and ColQwen	60
Table 14 qwen2.5 advantages	66
Table 15 Configuration of Qwen2.5-VL.....	68
Table 16 Performance of Qwen2.5-VL and State-of-the-art..	69
Table 17 Performance on pure text tasks of the 70B+ Instruct models and Qwen2.5-VL	70
Table 18 Performance of Qwen2.5-VL and other models on OCR, chart, and document understanding	70
Table 19 Open-domainDocVQAevaluationresultsonM3DOCVQA	71
Table 20 Closed-domainDocVQA evaluation results on MMLongBench-Doc	71
Table 21 Closed-domainDocVQAevaluation results onMP.....	72
Table 22 Comparison of different multimodal LMs within.....	72

Table 23 Overview of DocVQA Datasets Considered.....	74
Table 24 Answer-generationpilotresultsforQwen-2.5-VL-3Bon1k-questionsubsets.....	84
Table 25 Answer-generationpilotresultsforInternVL3-2Bon1k-questionsubsets.....	84
Table 26 RetrievalpilotresultsforColQwen(1Bvision–textencoder,768-dembeddings,HNSW).....	85
Table 27 RetrievalpilotresultsforColPaLI(PaLI-5Bencoder,1024-dembeddings, IVF-PQ).....	85

List of Abbreviations

ANLS: Average Normalized Levenshtein Similarity

ANN: Approximate Nearest Neighbor

BM25: Best Match 25 (a classical information retrieval algorithm)

BoW: Bag-of-Words

CBOW: Continuous Bag-of-Words

CNNs: Convolutional Neural Networks

DocVQA: Document Visual Question Answering

DPR: Dense Passage Retriever

FP32/FP16: 32-bit/16-bit Floating-Point

GUI: Graphical User Interface

HNSW: Hierarchical Navigable Small World

INT8/INT4: 8-bit/4-bit Integer

IR: Information Retrieval

LLMs: Large Language Models

LoRA: Low-Rank Adaptation

OCR: Optical Character Recognition

PEFT: Parameter-Efficient Fine-tuning

PTQ: Post-Training Quantization

QAT: Quantization-Aware Training

RAG: Retrieval-Augmented Generation

TF-IDF: Term Frequency-Inverse Document Frequency

UML: Unified Modelling Language

ViT: Vision Transformer

VLMs: Vision-Language Models

VRAM: Video RAM

List of Equations

Equation 1 corpus of documents	48
Equation 2 global set of page images	48
Equation 3 number of page	48
Equation 4 MaxSim score.....	50
Equation 5 Top-K page selection	50
Equation 6 Question answering output.....	51
Equation 7 Late Interaction	63
Equation 8 Contrastive Loss.....	63

Glossary

- **Cross-Encoders:** A type of interaction-based model where the query and document are processed together from the beginning, allowing for deep, token-level interaction between them to determine relevance.
- **Dense Retrieval:** A retrieval paradigm, also known as a Two-Tower Model, where the query and the document are encoded into fixed-size numerical vectors (embeddings) independently before their similarity is calculated.
- **Document Visual Question Answering (DocVQA):** An AI task that involves providing a model with a complex document (such as a PDF) and a natural language question, and expecting a relevant answer based on both the text and visual layout.
- **Embeddings:** A learned, dense, and relatively low-dimensional vector representation of a piece of data (like text or an image), where semantic similarity corresponds to geometric proximity in the vector space.
- **Few-shot Learning:** An extension of one-shot learning where a model is provided with a small number of examples (typically 2-32) in its prompt to better discern the pattern of a task.
- **Full Fine-tuning:** A traditional approach where all the parameters (weights) of a pre-trained model are updated during training on a new, task-specific dataset.
- **Information Retrieval (IR):** A field of computer science concerned with finding relevant information from a large collection of unstructured data that satisfies a user's information need, typically expressed as a query.
- **Late Interaction:** A retrieval paradigm that seeks to combine the efficiency of dense retrieval with the effectiveness of cross-encoders by delaying the query-document interaction until the final layers of the model, allowing for fine-grained matching with some pre-computation.
- **Model Quantization:** An optimization technique used to reduce the computational and memory costs of a neural network by converting its parameters from high-precision data types (like FP32) to lower-precision ones (like INT8).
- **Parameter-Efficient Fine-tuning (PEFT):** A method to adapt large pre-trained models by freezing the vast majority of the model's parameters and only training a small number of new or additional parameters.
- **Retrieval-Augmented Generation (RAG):** A hybrid architecture that enhances a language model by first retrieving relevant information from an external knowledge source and then providing that information as context to the model to generate a more accurate and factually grounded answer.
- **Sparse Retrieval:** Traditional information retrieval methods, such as TF-IDF and BM25, that operate primarily on lexical matching by analyzing the frequency and distribution of keywords from the query within documents.
- **Vision-Language Models (VLMs):** A class of multi-modal AI models that are designed to concurrently process and reason over information from both images and text, often by augmenting a large language model with a vision encoder.
- **Vision Transformer (ViT):** An architecture that adapts the Transformer model for computer vision by treating an image as a sequence of smaller, fixed-size patches. It uses self-attention to learn global relationships between all patches, making it effective for understanding overall layout.

- **Zero-shot Learning:** The capability of a model to perform a task without having seen any explicit examples of that task during its training, relying on the generalized knowledge acquired during pre-training to understand the task from a prompt alone

Chapter 1 : An Introduction

Visual Question Answering (VQA) is a well-established multi-modal AI task where a system is presented with an image and a natural language question related to that image, and is expected to generate an accurate answer. The traditional domain of VQA typically revolves around natural images, such as everyday scenes or photographs. For instance, a user might ask, “How many children are on the swings?” based on a park image. Answering such questions relies on capabilities like object detection, attribute recognition, and basic spatial reasoning.




(a) Q: What is the boy playing with?
A: teddy bear

(b) Q: Are there any animals swimming in
the pond? A: No

(c) Q: How many trees? A: 1

Figure 1 Example of question and generated answer for abstract images(VQA)

However, Document Visual Question Answering (DocVQA) represents a significantly more specialized and challenging subfield. While it shares the foundational goal of answering questions based on visual input, the nature of the “images” — complex documents like invoices, academic papers, contracts, and medical reports — introduces layers of complexity not present in standard VQA. The key challenge in DocVQA lies in the dense interplay between textual content and document layout, where meaning is embedded not only in text but also in structure, alignment, and visual formatting.



Q: What was the gross profit in the year 2009?

A: \$19,902

Figure 2 Example of question and generated answer for abstract images(DocVQA)

Three core factors distinguish DocVQA from conventional VQA:

- **Information Density:** Unlike natural images, documents contain a high concentration of semantic content. The answer may be a single word or number buried in a multi-page report, demanding precise extraction.
- **Layout Semantics:** The visual structure of documents—such as tables, forms, and headings—is crucial. For example, a number under a specific column in a table conveys a different meaning than the same number in a paragraph. Flattening documents through OCR often strips away this context, degrading comprehension.
- **Text-Centric Reasoning:** Questions in DocVQA often require interpreting text within a visual context, such as “What was the total amount due on the invoice dated May 15th?” rather than asking about visual features like colors or object counts.

With the exponential growth of digital data, organizations increasingly rely on automated systems to extract insights from such unstructured documents. The ability to query documents using natural language—asking questions like “What is the patient’s diagnosis?” or “What clause outlines the penalty terms?”—is a central goal in the broader domain of Document AI. However, building robust and scalable DocVQA systems remains a formidable challenge.

Traditional VQA models, designed for natural images, fall short when applied to documents due to their lack of layout understanding. Similarly, purely text-based approaches that depend on OCR and Retrieval-Augmented Generation (RAG) pipelines often ignore spatial structure, leading to loss of critical context. This is especially problematic in enterprise-scale documents, where key information is dispersed across dozens of pages, and the reading order or formatting carries semantic significance. Moreover, while Large Language Models (LLMs) are powerful, their input size limitations make them ill-suited for long documents, and they often hallucinate when overwhelmed.

On the other hand, Vision-Language Models (VLMs) offer better understanding of visual and layout features but are computationally expensive and struggle to scale efficiently. This creates a trade-off in current approaches: scalable retrieval systems often lack visual reasoning, while vision-aware systems sacrifice scalability.

To bridge this gap, this research proposes a **hybrid and flexible DocVQA framework** that combines the strengths of both paradigms. The core idea is a two-stage process:

1. **Page-level retrieval** using both visual and textual features to select relevant portions of the document.
2. **Answer generation** via a multi-modal model operating on the retrieved context.

This design enables efficient and accurate DocVQA over long, complex documents while remaining adaptable to different model backbones and retrieval strategies. The proposed system demonstrates notable improvements over traditional baselines across diverse document types and question categories. Ultimately, this framework lays the foundation for developing advanced AI assistants capable of deep, interactive understanding of visually rich documents.

.

1.1 Overview

Automatically understanding and extracting information from documents like invoices, legal contracts, and medical reports is a critical need for modern organizations. The field of Document AI is dedicated to this challenge, building systems to interpret these complex files. A major goal within this field is **Document Visual Question Answering (DocVQA)**, which aims to allow users to have a conversation with their documents by asking questions in plain, natural language.

However, achieving this is far more difficult than it sounds. The unique structure of documents, where meaning comes from a mix of text, tables, charts, and overall layout, presents a significant roadblock for even the most advanced AI models. This report will first break down these specific technical challenges in detail, and then propose a novel and more effective solution to improve how AI systems can understand documents at a deeper level.

1.2 Problem Statement

Visual Question Answering (VQA) has shown remarkable progress in handling natural images; however, applying the same methodologies to documents (DocVQA) reveals a significantly more complex challenge. Document images, particularly in enterprise or medical domains, often span multiple pages and exhibit dense, highly structured layouts involving text, tables, charts, and diagrams. Two critical issues hinder the effectiveness of existing VQA systems in this setting:

1. **Context Scale:** Enterprise documents are long and multi-page, often exceeding the input limits of large language models (LLMs). This limitation results in **factual hallucinations** and an inability to answer questions that require synthesizing information scattered across multiple pages.
2. **Visual Complexity:** OCR-based approaches used to extract text for LLMs typically strip away essential visual information—such as table structure, spatial relationships, or font styling—leading to significant information loss. As a result, models become **blind to layout and visual cues**, making it difficult to answer questions that depend on them.




Figure 3 Comparison of multi-modal document understanding pipelines. Previous works focus on (a) Single-page DocVQA that cannot

To address these challenges, researchers have explored several solutions. However, each approach presents its own **trade-offs**, leading to a persistent gap in building a robust and scalable DocVQA system.

Limitations of LLMs in DocVQA Systems

Large Language Models (LLMs) such as GPT-4, LLaMA 3, and PaLM have revolutionized natural language processing by demonstrating an impressive ability to perform a wide array of language-centric tasks without explicit task-specific supervision. This is largely due to their training on extensive textual corpora and their reliance on the Transformer architecture, which enables them to model long-range dependencies and nuanced semantic relationships in text.

What LLMs Do Well:

- Textual understanding and generation (e.g., answering questions, summarization).
- Reasoning based on linguistic cues.
- Transfer learning across domains when text is sufficient.

What LLMs Lack in DocVQA:

However, LLMs are **inherently unimodal**. This means they process only *textual tokens* and have no inherent understanding of:

- The **spatial layout** of text (e.g., whether a word is part of a header, footnote, or table cell).
- **Visual context**, such as lines, boxes, images, or charts embedded in the document.
- **Visual structure**, such as multi-column formats, nested tables, or hierarchical diagrams.

For example, when an LLM is fed OCR-extracted text from a table, it cannot infer *which cell belongs to which row or column*, because that information is purely visual and lost during linear text extraction.

Understanding large language models (LLMs) deployed by ChatGPT



This slide provides information regarding functioning of large language model which ChatGPT deploys. Such models are trained on huge amounts of textual data that revert in natural language text as response.




Figure 4 Large Language models architecture

This creates a **semantic gap** between the original visual document and what the LLM sees — a flat, structureless sequence of words. As a result, LLMs are not just blind to visuals, but also **deaf to layout semantics**, which are critical for understanding complex documents.

Emergence and Advantages of VLMs

To address the visual blindness of LLMs, **Vision-Language Models (VLMs)** have emerged. These models integrate visual perception capabilities into the language modeling pipeline, creating a **multimodal architecture**.

How VLMs Work:

- A **vision encoder** (e.g., ViT or ResNet) converts images (e.g., document pages) into dense embeddings.
- These embeddings are fused with **text embeddings** (from the document or question) and processed jointly.
- The **language decoder** (like in LLaVA or Qwen-VL) reasons over the combined image-text space and produces a response.

This design allows VLMs to “see” **both the visual layout** (where elements appear) and the **textual content** (what the elements say).

What VLMs Excel At:

- Understanding **tables, charts, and visual structures** that LLMs can't see.
- Answering questions about **spatial relationships** (“What is the value in the cell next to 'Total Revenue'?”).
- Performing **layout-aware reasoning**, such as identifying headers, footnotes, or segregating sections based on their position.

Limitations of VLMs in DocVQA

Despite these advances, VLMs still face **critical limitations**, especially when applied to large-scale document processing:

1. Context Window Limitations:

Just like LLMs, many VLMs are limited in how many tokens or image patches they can process at once. A long enterprise document (50+ pages) cannot fit into a single forward pass.

2. Computational Cost:

Processing high-resolution images of document pages is **GPU-intensive**, especially when dealing with full documents instead of cropped snippets.

3. Lack of Cross-Page Reasoning:

VLMs are generally **page-bound**. They might be able to answer questions *within a single page*, but struggle when answers are **distributed across multiple pages**.

4. No Dynamic Knowledge Integration:

VLMs do not *retrieve* relevant documents—they rely entirely on what’s given as input. So if a page containing the answer is not included in the input image batch, the model cannot find it.




Figure 5 Vision Language models architecture

Table 1 Trade-Off: LLMs vs. VLMs in DocVQA

CAPABILITY	LLMS (E.G., GPT, LLAMA)	VLMs (E.G., GPT-4V, QWEN-VL)
MODALITY	Text-only	Text + Visual (Multimodal)
VISUAL LAYOUT UNDERSTANDING	Absent	Present
CONTEXT CAPACITY	Limited	Also limited (esp. image resolution + length)
COMPUTATIONAL COST	Moderate	High
CROSS-DOCUMENT REASONING	Weak (due to linear OCR text)	Typically limited to one page at a time
SCALABILITY TO ENTERPRISE DOCS	with RAG	Poor due to image processing bottlenecks
EXPLAINABILITY / TRACEABILITY	Limited without retrieval	Not guaranteed

Thus, neither LLMs nor VLMs alone provide a complete solution:

- **LLMs** scale well with text and work with retrieval, but are visually blind.
- **VLMs** see the layout and page design, but choke on scale and lack dynamic memor

Attempts to Solve the Problem

Naive Fine-Tuning of VLMs: Why It Falls Short

A straightforward approach is to fine-tune a pre-trained **Vision-Language Model (VLM)**, such as LLaVA, IDEFICS, or Qwen-VL, directly on DocVQA datasets. VLMs are powerful multi-modal models that process both visual and textual information. They integrate a **vision encoder** (e.g., ViT) with a language decoder, allowing them to jointly reason over images and text. Fine-tuning such models can, in theory, help adapt them to document-specific formats, layouts, and question types.

However, full fine-tuning of large VLMs introduces several critical limitations:

- **Computational Expense:** Updating billions of parameters requires immense computational resources and high memory GPUs, making it inaccessible for many real-world applications.
- **Catastrophic Forgetting:** Fine-tuning the entire model can degrade the general-purpose capabilities learned during pretraining, making the model overly specialized to the fine-tuning dataset.
- **Static Knowledge:** Once the model is fine-tuned on a specific dataset, it cannot adapt to new documents or domains without re-training.
- **Poor Explainability:** Even after fine-tuning, the model generates answers without clear grounding in specific content, which reduces trust and interpretability.

To partially mitigate these challenges, **Parameter-Efficient Fine-Tuning (PEFT)** techniques like **Low-Rank Adaptation (LoRA)** have been introduced.

LoRA: A Partial Fix with Remaining Flaws

LoRA significantly reduces training costs by injecting small, trainable parameter matrices into the model while freezing the rest. This allows adapting large models with fewer resources.

Pros of LoRA:

- Lower memory and compute requirements.
- Maintains most pre-trained knowledge while allowing domain-specific adaptation.

Still, LoRA Doesn't Solve Core Issues:

- **Static Knowledge:** The model still cannot adapt to new or evolving documents without retraining.
- **Hallucinations Persist:** LoRA does not inherently improve the factual grounding of answers.
- **Lack of Transparency:** Like full fine-tuning, it still fails to show users where the answer came from in the document.

Thus, even with LoRA, fine-tuning remains **inflexible, static, and hard to trust** in document-rich environments.




Figure 6 Our reparameterization. We only train A and B

Limitations of Traditional RAG in Document Visual Question Answering (DocVQA)

Retrieval-Augmented Generation (RAG) has emerged as a powerful paradigm for enhancing language models with access to external knowledge. In this setup, a retriever first selects the most relevant documents (or document chunks) from a large corpus based on a given query. These retrieved texts are then passed to a generator (typically an LLM), which synthesizes an answer based on both the query and the retrieved content.

This framework brings multiple advantages:

- **Improved factual accuracy**, by grounding answers in real data.
- **Scalability**, since the model doesn't need to store all knowledge internally.
- **Interpretability**, as users can trace answers back to retrieved contexts.

RAG Architecture Model




Figure 7 RAG architecture

But Here's the Catch:

Traditional RAG systems are **text-centric by design**. Both the retriever and generator modules operate exclusively in the **textual modality**:

- **Retrievers** (like BM25, DensePassage, or ColBERT) match queries to preprocessed text chunks (e.g., paragraphs).
- **Generators** (like GPT or T5) consume only text to produce answers.

This design is effective in scenarios like question answering over Wikipedia or scientific papers where all relevant information is **well-structured and fully represented in text**.

However, when applied to **document question answering**, especially over **visually complex documents** (e.g., invoices, forms, multi-column reports), this framework **collapses** due to a crucial limitation:

Traditional RAG is blind to visuals.

The retriever cannot “see” images, layouts, or spatial relationships.

The generator cannot “reason” over tables, figures, or embedded charts.

Even if OCR is used to extract the textual content of each page, the resulting plain text **flattens** the document structure and **discards visual semantics**, leading to:

- Loss of **table-cell alignment**
- Disconnection between **figures and captions**
- Inability to interpret **multi-column layouts**
- Misunderstanding of **text hierarchies** (headers vs. footnotes)

Thus, traditional RAG becomes **fundamentally insufficient** in scenarios where:

- The answer relies on **how** information is presented (not just what is said),

- Or when the query refers to **visual cues**, like “the value in the table below” or “the image on the top right.”

The Missing Link: Can RAG Work with Images?

Despite its limitations, RAG offers something **hugely valuable**: the ability to scale document QA to large corpora. It can retrieve the **relevant pages**, **cross-reference multiple sources**, and provide a **flexible memory** that helps overcome context-window limitations in LLMs or VLMs.

So, if we could somehow **enable RAG to retrieve and reason over visual content**, we would have the best of both worlds:

Scalable retrieval across thousands of pages

Accurate reasoning that incorporates layout and visual elements

Table 2 Limitations of Traditional RAG for Document Visual Question Answering

COMPONENT	TRADITIONAL RAG	REQUIREMENT IN DOCVQA	LIMITATION
RETRIEVER	Based on textual embeddings (e.g., BERT)	Multimodal retrieval using both text and layout	Unable to retrieve based on visual or spatial cues
GENERATOR	Text-only language model (e.g., GPT, T5)	Ability to reason over images and structured layouts	Lacks visual reasoning capabilities
INPUT REPRESENTATION	Flattened OCR text	Structured documents with rich visual layout	Loses hierarchy, spatial alignment, and semantics
ANSWERING SCOPE	General question answering over plain text	Deep understanding of document semantics and layout	Insufficient for layout-dependent queries

The Core Trade-Off

Ultimately, these attempts present a critical trade-off:

Table 3 Core Trade-Offs in Current DocVQA Approaches

APPROACH	HANDLES LONG CONTEXT	PRESERVES VISUAL LAYOUT	DYNAMIC KNOWLEDGE	TRANSPARENT REASONING
LLM + OCR	✓	✗	✗	✗
VLM FINE-TUNING	✗	✓	✗	✗
LORA	✗	✓	✗	✗
RAG (TRADITIONAL)	✓	✗ (OCR-only)	✓	✓

This trade-off defines the **core research challenge**:

Can we combine the **context scalability of RAG** with the **layout-awareness of VLMs** into a single framework that supports dynamic, transparent, and accurate reasoning over complex multi-page documents?

This work aims to answer that question by designing and implementing a **DocVQA pipeline** that tightly integrates **multi-modal retrieval** with **vision-language reasoning**, laying the foundation for more robust document AI systems.




Figure 8 RAG VS Fine Tuning

Review of Information Retrieval Techniques and the Persistent Visual Gap

The effectiveness of any RAG framework hinges on the quality of its **retrieval** component. A robust retriever must surface the most relevant document chunks for a given query;

otherwise, no amount of downstream reasoning can compensate for a missing or irrelevant context. Over time, Information Retrieval (IR) has evolved through two main paradigms—classical (sparse) retrieval and neural retrieval—each bringing improvements in text-based matching but neither addressing the challenge of visual information in documents.

1. Classical Information Retrieval (Sparse Retrieval)

Classical IR methods rely on lexical matching between query terms and document text:

- **Key Algorithms**
 - **TF-IDF** weights terms by how frequently they appear in a document versus the corpus, highlighting document-specific keywords.
 - **BM25** refines TF-IDF with probabilistic considerations, forming a strong baseline for keyword search.
- **Limitations**
 - **Semantic Gap**: These methods match words literally, failing to handle synonyms or word sense (“automobile” vs. “car”).
 - **Modality Restriction**: They process text only, with no mechanism to incorporate images, layout cues, or spatial relationships.




Figure 9 Information Retrieval

2. Neural Information Retrieval Architectures


Neural IR has introduced semantic understanding and scalable search, broadly categorized as:

1. **Representation-based Similarity (Dense Retrieval / Two-Tower Models)**
 - Query and document chunks are independently encoded into fixed-size vectors (e.g., via BERT).
 - Retrieval is a fast nearest-neighbor search over precomputed document embeddings.
 - **Pros**: Highly scalable; embeddings can be indexed for rapid lookup.
 - **Cons**: A single vector can bottleneck nuanced information and still ignores visual elements.
2. **Interaction-based Models (Cross-Encoders)**
 - Query and document text are concatenated and jointly processed, capturing fine-grained token-level interactions (e.g., BERT as a ranker).
 - **Pros**: State-of-the-art ranking effectiveness.

- **Cons:** Impractically slow for first-stage retrieval; operates on text only.

3. Late Interaction Models (e.g., ColBERT)

- Encodes each document token into its own vector and delays query–document interaction until the final layers.
- Balances precomputable embeddings with finer matching than dense retrieval.
- **Pros:** Better accuracy than single-vector models while remaining scalable.
- **Cons:** Increased storage and complexity; still strictly text-based.



Despite these advances in retrieving **textual** information—reducing semantic gaps, improving ranking, and scaling to massive corpora—**none of these IR approaches incorporate visual or layout features**. As a result, a RAG system built atop classical or neural IR remains fundamentally **blind to images, tables, and spatial structure**. Even with the most sophisticated text retriever, the RAG pipeline will fail to surface critical visual contexts—and thus cannot support questions whose answers depend on how information is visually arranged. This persistent modality gap underscores the need for a **truly multimodal retrieval** strategy before any downstream generation can effectively address DocVQA tasks.

In summary, despite significant advances in both language-only and vision-language models, as well as sophisticated retrieval techniques, current DocVQA systems remain impeded by a persistent duality: they can either scale to long, multi-page documents at the cost of all visual structure, or they can exploit rich layout information but fail to handle enterprise-scale contexts. Traditional fine-tuning—even when parameter-efficient—yields static, opaque models that cannot adapt dynamically, while standard RAG pipelines, grounded in text retrieval, remain blind to tables, charts, and spatial hierarchy. Even the most advanced neural IR architectures, from dense retrieval to late-interaction re-rankers, operate solely on text embeddings and thus cannot surface visually encoded semantics. This convergence of limitations underscores the urgent need for a truly multimodal retrieval and reasoning framework—one that unites scalable, context-aware search with layout-sensitive visual understanding—to bridge the gap and enable robust, explainable question answering over complex, real-world documents.

1.3 Scope and Objectives

Scope of Work

The aim of this project is to design, build, and evaluate a domain-agnostic Document Visual Question Answering (DocVQA) system. This system will leverage a novel Retrieval-Augmented Generation (RAG) pipeline that uses a Vision-Language Model (VLM) to interpret complex, multi-page PDF documents. By embedding and retrieving entire page images, the system will bypass brittle and slow OCR-based text extraction processes, enabling it to natively understand the interplay of text, layout, and visual elements like tables and charts. The ultimate goal is to deliver a scalable and efficient solution that provides accurate, factually grounded, and traceable answers in near real-time.

To maintain a clear research focus and ensure feasibility within the given timeframe, the scope of this project is defined as follows:

In-Scope:

- **Benchmark Datasets:** The system will be evaluated using a diverse set of standard benchmarks to comprehensively assess its capabilities. This includes datasets focused on long-context reasoning (*MPDocVQA*, *MMDocLongBench*, *ArxivQA*) and complex visual layout understanding (*DUDE*, *TAT-DQA*, *SlideVQA*).
- **Language:** The project is limited to English-language documents and questions, aligning with the language distribution in the selected benchmarks.
- **Text Modality:** While the architecture is theoretically capable of processing both printed and handwritten text, primary evaluation will be conducted on datasets that predominantly contain printed text.
- **Model Foundation:** The project will leverage **pre-trained, open-source** Vision-Language Models (VLMs) as a foundational component.
- **Model Fine-tuning:** The project **will involve fine-tuning** the pre-trained VLMs on task-specific datasets to optimize their performance for document understanding.

Project Objectives

This research project aims to achieve the following core objectives:

1. **Design a Visually-Aware and Performant Retrieval Mechanism:** To develop and optimize a retrieval module that leverages both textual and visual cues to accurately

extract relevant document segments. The design will prioritize low-latency performance to enable near real-time interaction.

2. **Build and Fine-Tune an End-to-End DocVQA System for Deep Comprehension:** To create a cohesive RAG pipeline that integrates the retrieval module with a Vision-Language Model (VLM). The system will be fine-tuned to achieve deep multimodal comprehension, ensuring it can accurately interpret tables, charts, and complex layouts.
3. **Ensure High Factual Accuracy and Verifiability:** To implement and refine the system to mitigate model hallucinations by grounding all generated answers in the retrieved source content. The system must produce verifiable outputs with clear references to the source document.
4. **Establish Rigorous Baseline Comparisons:** To implement and evaluate strong baseline models, including traditional text-based RAG pipelines. This will serve to quantitatively measure the added value and superior performance of the proposed visually-aware, OCR-free approach.
5. **Demonstrate State-of-the-Art Performance via Comprehensive Evaluation:** To conduct a robust quantitative evaluation on standard benchmarks (such as MP-DocVQA, MMLongBench-Doc, and TAT-DQA). The goal is to demonstrate that the system surpasses existing baselines and achieves state-of-the-art results, particularly on tasks requiring long-context and visual reasoning.
6. **Architect for Scalability:** To design the solution's architecture to be capable of scaling efficiently to handle large corpora of documents and support a high volume of concurrent user queries.

1.4 Report Organization (Structure)

This report systematically addresses the challenges and solutions in the field of Document Visual Question Answering (DocVQA). The report is organized into five main chapters as follows:

- **Chapter 1: Introduction** provides an overview of the Document AI field and defines the core problem this research addresses: the limitations of current systems in handling long and visually complex documents. It also outlines the project's scope and objectives, as well as the methodology and work plan followed.
- **Chapter 2: Related Work (Literature Review)** reviews the theoretical background and foundational concepts, such as Large Language Models (LLMs), Vision-Language Models (VLMs), and Retrieval-Augmented Generation (RAG). It presents a comprehensive survey of previous research in the field, with a comparative analysis of the most prominent methodologies, identifying the research gaps this project aims to address.
- **Chapter 3: The Proposed Solution** offers a detailed description of the proposed architecture (M3DocRAG), which integrates multimodal information retrieval and question answering using advanced models. This chapter explains the solution's methodology step-by-step, from document embedding and page retrieval to answer generation, along with the rationale for selecting the models and tools used.
- **Chapter 4: Implementation, Experimental Setup, & Results** focuses on the practical aspects of the project. It details the system's implementation, the experimental setup used to evaluate the performance of the proposed solution, and presents the results obtained from experiments and tests on standard benchmark datasets.

- **Chapter 5: Discussion, Conclusions, and Future Work** discusses the results obtained, compares them to previous work, and presents a critical evaluation of the proposed solution's strengths and weaknesses. The report concludes with a comprehensive summary of the main contributions and suggests future research directions that can build upon this work

1.5 Work Methodology

The methodology for this project was designed to integrate **rigorous academic research** with **agile software development practices**, ensuring both **scientific validity** and **practical implementation efficiency**.

A **hybrid approach** was adopted, combining:

- An **Experimental Research Framework** for formal investigation.
- An **Agile-inspired development process** for iterative implementation.

1. Experimental Research Framework

Phase 1: Literature Review & Problem Definition

This phase laid the foundation for the project. A deep and structured survey of over **25+ research papers** was conducted, covering:

- DocVQA vs traditional VQA
- Limitations of existing approaches (e.g., context length, layout blindness)
- Architectures of LayoutLM, Donut, BLIP, Qwen-VL, etc.
- Retrieval techniques (e.g., MaxSim, ColPaLi)
- Long-context modeling strategies (FiD, HyDE, ColBERT)

A comparative matrix was maintained to track:

- Visual/textual input modalities
- Document layout handling
- Context window size
- Dataset support

Key Insight: No existing method achieved both **layout awareness** and **multi-page reasoning**, which led to our central hypothesis.

Phase 2: Hypothesis & System Design

Hypothesis: A visually-aware, page-level retrieval mechanism feeding into a strong VLM can significantly improve DocVQA performance on complex documents.

A **modular architecture** was proposed:

- Visual/textual encoders per document page
- A layout-sensitive retriever (e.g., FAISS + ColPaLi)

- A VLM answering model (Qwen-VL or similar)

Two experimental tracks were defined:

- **Track A** – Zero-/Few-shot (no fine-tuning)
- **Track B** – Fine-tuned using parameter-efficient methods (e.g., LoRA)

Phase 3: System Implementation

- **Data pipeline:** image resizing, visual layout parsing
- **Retriever module:** Embedding extraction, FAISS-based indexing, MaxSim scoring
- **VLM integration:** Prompting (Track A) vs fine-tuning (Track B)
- Unified input formatting for consistency across tracks

Phase 4: Evaluation

The system was benchmarked on:

- **Long-context datasets:** mpdocvqa, mmlongbenchdoc, arxivqa
- **Layout-heavy datasets:** dude, tatdqa, slidevqa

Both tracks were tested under identical settings.

Evaluation Metrics:

- Accuracy
- F1 Score
- ANLS (Average Normalized Levenshtein Similarity)

Phase 5: Analysis & Conclusion

- Performance comparison across tracks
- Qualitative error analysis
- Practical trade-offs between zero-shot vs fine-tuned
- Limitations (e.g., model size, retrieval errors)
- Proposed directions (e.g., joint retriever-reader training)

2. Agile-Inspired Development Process

While the research was grounded in formal experimentation, the development of the system followed an **Agile-inspired workflow**, allowing iterative progress and continuous refinement. The work was broken down into focused sprints, each lasting approximately two weeks, with the following structure:

SPRINT	OBJECTIVE
---------------	------------------

SPRINT 0	Read and analyze 25+ papers, define hypothesis, identify benchmarks & metrics
SPRINT 1	Build data loaders, process OCR & layout, resize images, create unified inputs
SPRINT 2	Implement page-level retriever using FAISS & embeddings, test retrieval quality
SPRINT 3	Implement Track A (zero-shot VLM pipeline) with proper prompts
SPRINT 4	Fine-tune VLM (Track B) using LoRA on task-specific data
SPRINT 5	Evaluate both tracks, compute metrics, perform error analysis, document results

Table 4 Agile Process

3. SMART Goals Framework

To ensure the project objectives are actionable, results-driven, and time-conscious, the SMART framework was adopted as a guiding methodology for planning and evaluating progress throughout the development lifecycle. Each goal was designed to meet the following five criteria:

1. **Specific:**
Every objective is clearly defined, leaving no room for ambiguity. For instance, instead of stating general intentions like “improve the model,” the project aims to “develop a visual retrieval component that selects the most relevant document pages based on layout-aware embeddings.”
2. **Measurable:**
Goals are paired with measurable outcomes, allowing for clear tracking of progress. Examples include performance benchmarks such as achieving a 5% increase in ANLS (Average Normalized Levenshtein Similarity) or Recall@K when compared to baseline methods.
3. **Achievable:**
All defined goals are realistic within the project’s scope, technical capabilities, and available resources. Instead of targeting full retraining of large-scale models, the project focuses on using pre-trained vision-language models with parameter-efficient fine-tuning techniques (e.g., LoRA).
4. **Relevant:**
Each objective is directly aligned with the core research challenge: enhancing the performance of Document Visual Question Answering (DocVQA) systems, especially in long-context and layout-sensitive scenarios. Irrelevant or peripheral features (e.g., mobile deployment, multilinguality) were intentionally excluded.
5. **Time-Bound:**
Goals are tied to specific timelines and milestones within the project’s Gantt chart. For example:
 - Literature review and problem definition were allocated two months (March–April).
 - Retrieval module development was planned for Sprint 2 (mid-May).
 - Fine-tuning and experimental evaluation spanned Sprints 3–5 (June).

This structured approach ensures that the project not only remains aligned with its research objectives, but also progresses in a manageable, trackable, and outcome-oriented manner.



Figure 11 SMART Business Goals

GOAL	SMART CRITERIA ALIGNMENT
IMPROVE DOCVQA PERFORMANCE BY 5% OVER VLM BASELINE	Measurable & Specific
BUILD PIPELINE IN 24 WEEKS WITH 2-WEEK SPRINTS	Time-bound & Achievable
SUPPORT AT LEAST 3 DIVERSE BENCHMARKS USE OPEN-SOURCE MODELS (QWEN-VL, COLPALI, ETC.)	Relevant & Specific Achievable & Transparent
MAINTAIN REPRODUCIBILITY THROUGH MODULAR DESIGN	Measurable & Relevant

Table 5 SMART Business Goals

1.6 Work Plan (Gantt chart)

The implementation of the project followed a structured work plan spanning approximately 24 weeks, divided into logically sequenced phases that align with both the experimental research methodology and Agile development principles. Below is an overview of the project timeline and key activities:

Week(s)	Phase	Activities
1	<i>Literature Survey & Planning</i>	Conducted an extensive review of 25+ papers on DocVQA, retrieval-augmented models, and Vision-Language Models. Formulated the research hypothesis and finalized benchmark datasets.
6–10	<i>Data Preparation</i>	Downloaded and explored selected datasets (mpdocvqa, mmlongbenchdoc, tatdqa, slidevqa, dude). Built pipelines for OCR extraction, page-image resizing, and unified data formatting.

10–12	<i>Retriever Implementation</i>	Implemented a page-level retrieval system using visual and textual embeddings. Performed ablation testing on retriever accuracy (e.g., recall@k, relevancy scores).
12–14	<i>Track A – Zero-shot VLM Pipeline</i>	Built and tested a baseline DocVQA pipeline using frozen VLMs without fine-tuning. Designed prompts and validated on a subset of questions.
14–16	<i>Track B – Fine-tuning VLMs</i>	Applied parameter-efficient fine-tuning (LoRA, QLoRA) on selected VLMs using task-specific documents. Evaluated improvements over the baseline.
16–18	<i>System Integration</i>	Integrated retriever with the reader (VLM) into a single end-to-end RAG-style pipeline. Handled input formatting and modality alignment.
18–20	<i>Evaluation & Analysis</i>	Ran full-scale evaluation across all selected datasets. Measured Accuracy, F1, and ANLS. Performed error analysis and comparison with baseline results.
20–24	<i>Documentation & Finalization</i>	Finalized all results, wrote thesis chapters, generated visualizations, and prepared project presentation.

Table 6 Work Plan

The plan was followed using Agile-inspired sprints, allowing quick adaptations. Sprint reviews guided architecture changes, model selection, and data handling strategies.




Figure 12 Gantt chart

Chapter 2: Related Work (Literature Review)

2.1 Background

The Evolution of Document Question Answering (Document QA)

Document question answering (**QA**) has evolved from simple **OCR**-and-text pipelines to sophisticated vision-language systems. Early work such as the **DocVQA** dataset demonstrated that answering questions on document images requires not only reading text but understanding the layout and structure of the page [Source].

1. Limitations of Traditional OCR-Based Approaches

Traditional pipelines first apply **OCR** or PDF parsing to extract text, then use text-based models (e.g. **BERT** or **T5**) to answer questions. However, many documents contain critical information in visual form – tables, charts, figures, and complex layouts – that are lost or degraded by **OCR**. Mathew et al. note that performance on **DocVQA** is particularly poor when “understanding the structure of the document” is crucial [Source].

2. The Rise of Layout-Aware and Vision-Language Models (VLMs)

To address this, the field began incorporating visual features and layout cues. Models like **LayoutLM** and its successors [Source], **DONUT**, etc., explicitly model the spatial arrangement of text and image regions. In parallel, large pre-trained vision-language models (**VLMs**) such as DeepMind’s **Flamingo** were developed to seamlessly integrate image and text inputs [Source]. **Flamingo** and similar **VLMs** “bridge powerful pretrained vision-only and language-only models” and can ingest arbitrarily interleaved visual and textual data [Source].

Document-specialized models have built on these ideas: for instance, **DocPrompt** emphasizes that high-capability document **QA** requires learning textual content, layout, and image priors jointly [Source], and **DocVLM** augments a base **VLM** with an **OCR** encoder to capture text+layout and compress it into learned query tokens [Source].

3. The Shift to Retrieval-Augmented Generation (RAG) for Documents

As vision-enabled models emerged, there has also been a shift toward retrieval-augmented **QA** architectures for documents. Classic **RAG** systems retrieve relevant text passages from a corpus of documents to ground a language model. But treating documents purely as text causes loss of visual semantics.

Recent work therefore performs retrieval over image embeddings or hybrid representations. For example:

- **VisRAG** embeds each document page as a whole image using a vision encoder, and then retrieves those image embeddings to feed into a generative **VLM** [Source]. By avoiding the **OCR** step, **VisRAG** “maximizes the retention and utilization of the data information in the original documents” and eliminates parsing-induced information loss [Source].
- **VDocRAG** frames all documents (PDFs, PPTs, etc.) as images and learns a retrieval encoder with self-supervised vision-language objectives. Tanaka et al. show that **VDocRAG**’s unified image-based retrieval “outperforms conventional text-based **RAG** and has strong generalization capability” on an open-domain document **QA** benchmark [Source].

In these systems, retrieval is typically handled with dense encoders or late-interaction models: e.g. dual-encoders that map queries and document images into a vector space, often with refinements such as multi-granular indexing (retrieving by document page, paragraph, or chart as needed).

4. Handling Multi-Page and Long-Context Documents

A new generation of “multi-modal **RAG**” frameworks extends this idea to long or multi-page documents.

- **M3DocRAG** (Cho et al., 2024) argues that many real queries require pooling evidence across pages or even multiple documents, and that visual elements (charts, figures) must be included in the retrieval [Source]. **M3DocRAG** uses a multi-modal retriever to find relevant pages (not just lines of text) and then conditions a multi-modal language model on those pages to answer the question, yielding state-of-the-art results [Source].
- **Hi-VT5**: Introduced by Rubèn Tito et al. for the Multi-Page DocVQA task, this is a hierarchical T5-based model whose encoder first compresses each page into a summary, and whose decoder then generates the answer from these page summaries [Source].
- **DocVLM** (Wu et al., 2024) shows that one can preprocess **OCR** from many pages into a compact set of learned queries for a **VLM**, enabling the model to handle more pages in its context window [Source].

5. The Trend Towards General-Purpose and Domain-Agnostic Systems

Across all these efforts, a common trend is moving toward general-purpose, domain-agnostic systems. Models are now often trained or evaluated on collections of diverse documents rather than narrow domains. For instance, **VDocRAG** introduces **OpenDocVQA**, a unified multi-domain benchmark spanning charts, forms, reports, and slides [Source]. **M3DocRAG** explicitly targets both closed-domain (specific corpus) and open-domain retrieval settings [Source]. Similarly, **DocPrompt** is motivated by industrial needs to apply **QA** to arbitrary business documents with minimal annotation, focusing on strong zero-/few-shot generalization [Source].

Summary and Conclusion

In practice, contemporary systems combine dual encoders (separate vision/text encoding pipelines), fusion modules (cross-attention layers or learned adapters), and hierarchical or iterative reasoning to handle multi-page inputs. They leverage the latest **VLMs** (e.g. **GPT-4o**, **Qwen-VL**, **LLaVA**) as backbone generators, augmenting them with visual retrieval and chunking mechanisms.

In summary, document **QA** has progressed from **OCR**-only pipelines to integrated vision-language **RAG** frameworks that preserve visual content (layout, graphics, tables) and retrieve information in a granular, multimodal fashion [Source].

Table 7 Key Techniques in Document QA

CATEGORY	EXAMPLES	STRENGTHS	LIMITATIONS
----------	----------	-----------	-------------

1. OCR + TEXT MODELS	BERT, T5 on DocVQA	Simple, fast, well-understood	Lose layout & visual info (charts, tables, structure)
2. LAYOUT-AWARE MODELS	LayoutLM, DONUT	Understand layout, spatial info	Still limited context size, not scalable to long docs
3. VISION-LANGUAGE MODELS	Flamingo, DocVLM	Jointly process images & text, keep visual features	Expensive, may not scale to many pages
4. TEXT RAG	Classic RAG (text retrievers)	Dynamic, low-cost, traceable answers	Visual content lost during OCR
5. VISUAL/MULTIMODAL RAG	VisRAG, VDocRAG, M3DocRAG	Retrieves whole images/pages, preserves layout & visuals	High latency, complex to train
6. HIERARCHICAL MODELS	Hi-VT5, DocVLM (2024)	Summarize multi-page docs, reduce context overload	Still require good preprocessing or summarization

Summary Points:

- **OCR-based methods** are fast but lose layout.
- **Layout-aware models** add structure but don't scale well.
- **VLMs** preserve both text and visuals but are resource-heavy.
- **RAG methods** bring dynamic context and traceability.
- **Multimodal RAG (like M3DocRAG)** = the current state-of-the-art, combining the best of both worlds.

2.2 Literature Survey

This section presents an overview of key studies in Retrieval-Augmented Generation (RAG), covering both foundational text-based models and their extensions to Document Visual Question Answering (DocVQA).

Lewis et al. (2020) – Retrieval-Augmented Generation (RAG)

- Introduced the RAG framework for open-domain question answering, combining dense retrieval (via DPR) with sequence generation.

- Achieved state-of-the-art results on multiple benchmarks (Natural Questions, WebQuestions, CuratedTrec, TriviaQA).
 - Demonstrated that RAG can generate accurate responses even when the exact answer isn't present in retrieved passages.
 - Provided strong evidence of RAG's generality across knowledge-intensive NLP tasks.
-

Muludi et al. (2024) – Document QA with Text-based RAG

- Applied standard text-based RAG (using GPT-3.5-turbo and FAISS retrieval) to document question answering.
 - Processed OCR-extracted text chunks from PDFs and generated answers using a traditional RAG pipeline.
 - Achieved high accuracy ($F1 \approx 0.88$), confirming the viability of RAG on document tasks.
 - However, performance depends heavily on OCR quality, and lacks multimodal reasoning.
-

Adjali et al. (2024) – Multi-Level RAG for VQA

- Proposed a multi-stage RAG model combining entity-level and passage-level retrieval in a joint training setup.
 - Demonstrated improved retrieval relevance and answer accuracy on the VIQuAE knowledge-based VQA benchmark.
 - Remains limited to textual inputs.
-

Long et al. (2025) – ReAuSE: Retrieval-Augmented Visual QA

- Proposed a multimodal, unified autoregressive retriever-generator for visual question answering.
 - Integrated retrieval within the generation loop and reported gains of 2.9–9.6 % on OKVQA and A-OKVQA.
 - Represents a core step toward end-to-end multimodal RAG systems.
-

M3DocRAG (2024)

- Developed for open-domain, multi-hop Q&A over multi-page PDFs.
- Uses ColPaLi-based page-level embeddings, MaxSim retrieval, and Qwen2-VL generation.
- Achieved leading results on M3DocVQA ($F1 = 36.5$), MMLongBench-Doc ($F1 = 22.6$), and MP-DocVQA.
- Key contribution: scalable multimodal retrieval with cross-page reasoning.

VisRAG (2024)

- Processes full-page images via vision-language models (e.g., BLIP-2), without OCR.
- Uses concatenated or weighted retrieval across page embeddings.
- Achieves 25–39 % improvement over text-based RAG in DocVQA and InfographicQA.
- Highlights the importance of preserving visual structure.

VisDoMRAG (2024–2025)

- Utilizes dual-mode retrieval from both visuals and text, followed by evidence curation.
- Integrates chain-of-thought reasoning for generating more coherent multimodal answers.
- Records 12–20 % gain on the VisDoMBench benchmark.
- Particularly effective for complex, multi-hop queries.

VDocRAG (2025)

- Incorporates self-supervised visual pre-training on structured elements like tables and charts.
- Introduces the OpenDocVQA dataset for multimodal document QA evaluation.
- Demonstrates superior performance to text-only RAG on visually structured content.

MagicLens (ICML 2024)

- Trains a dual-encoder model using triplets (image, instruction, image) for instruction-conditioned visual region retrieval.
- Marks a major shift toward semantic, fine-grained region retrieval within documents.
- Not yet evaluated specifically in the DocVQA setting.

PDF-WuKong (2024)

- Introduces sparse block-level sampling across long documents.
- Employs an end-to-end RAG pipeline including text and visual blocks.
- Achieves an 8.6 % F1 boost over baselines on the PaperPDF dataset.
- Balances retrieval granularity with processing efficiency.

Paper	Methodology	Input Modality	Key Contributions	Limitations
<i>Lewis et al. (2020) – RAG</i>	Classical RAG	Text only	Introduced RAG architecture combining retrieval and generation in open-domain QA.	Not applicable to multimodal inputs.
<i>Muludi et al. (2024)</i>	Classical RAG + OCR	OCR + text	Applied GPT-3.5 and FAISS retrieval to DocVQA using textual chunks.	OCR-dependent; lacks visual awareness.
<i>Adjali et al. (2024)</i>	Multi-level RAG	Text (structured entities)	Combines entity-level and passage-level retrieval in a joint retriever-generator setup.	No support for visual layout or images.
<i>Long et al. (2025) – ReAuSE</i>	Unified retriever + generator	Image + text	Integrates retrieval and generation using a multimodal autoregressive model.	Still limited in handling long documents.

<i>M3DocRAG</i> (2024)	Multimodal RAG	Document image + text	Introduced ColPaLi for multimodal embedding and MaxSim for page retrieval.	Retrieval is static and not question-aware.
<i>VisDoM</i> (2024)	Visual document retriever	Document image	Performs document-wide image embedding with joint retrieval.	Resource intensive; less scalable.
<i>VisRAG</i> (2023)	Patch-based RAG	Image patches + OCR	Uses VLT5 and patches to retrieve relevant document regions.	Fine-grained but lacks global context.
<i>MagicLens</i> (2024)	Instruction-tuned retriever	Image + text + instruction	Introduces self-supervised visual retriever guided by open-ended instructions.	Dependent on instruction formulation quality.
<i>PDF-WuKong</i> (2024)	Sparse multimodal retriever	Long PDF pages	Improves document-level retrieval via sparse sampling of visual features.	Requires complex training and annotation.

Table 8 papers comparison

2.3 Analysis of Related Work

VisRAG Recent vision-language QA systems based on retrieval-augmented generation (RAG) have explored different design choices for retrieving and leveraging document information. VisRAG pioneered a fully vision-centric pipeline, embedding whole document pages with a vision-language model (VLM) instead of extracting text. Its retriever is a dense bi-encoder that maps queries and page images into a shared embedding space. By processing pages as images, VisRAG avoids OCR-induced information loss and captures layout cues. This dense (image-based) retrieval strategy maximizes raw data retention, yielding significant gains (20–40%) over text-only RAG on multi-modal documents. However, fully vision-based retrieval can incur higher computational cost for large corpora and may overlook fine-grained textual semantics that OCR would capture.

M3DocRAG M3DocRAG introduces a multi-modal retrieval framework for multi-page, multi-hop QA. Unlike VisRAG’s image-only embeddings, M3DocRAG employs late-interaction retrievers (ColPaLi and ColQwen) that jointly encode page images and extracted text. Each document page is converted to an RGB image; visual embeddings and text tokens are indexed. At query time, a text query is compared against this multi-modal index using a MaxSim operator to retrieve relevant pages. This design preserves both visual layout (charts, figures) and textual context, enabling chart and table reasoning. Empirically, M3DocRAG achieves state-of-the-art performance on DocVQA tasks by flexibly switching between closed- and open-domain settings. Its modular retriever-generator architecture is explainable

and decoupled, though scaling to millions of pages requires approximate nearest-neighbor indexing and efficiency tuning.

VDocRAG VDocRAG extends RAG to visually rich documents—PDFs, slides, charts, and tables—by treating each page as an image and introducing vision-augmented pre-training for retrieval. A vision-augmented LLM compresses visual content into dense token embeddings aligned with text. The VDocRetriever finds similar page images in this joint space, and the VDocGenerator answers questions conditioned on retrieved images. By regarding charts and tables as first-class citizens, VDocRAG outperforms text-only RAG on multimodal datasets and generalizes strongly to new layouts. The downside is compute-intensive pre-training and dense indexing, which assume unified image-format inputs.

DocPrompt DocPrompt adopts a contrasting philosophy: instead of retrieval at inference, it relies on large-scale, continued pre-training of a document-understanding model. Building on layout-aware transformers (e.g., ERNIE-Layout), it ingests both text and layout tokens through weakly supervised data generation and multi-stage training. DocPrompt achieves state-of-the-art results on several DocQA benchmarks without external retrieval or OCR pipelines. However, it cannot explicitly ground answers in retrieved pages, and extending to cross-document reasoning remains challenging due to its monolithic architecture.

Lewis et al. (2020) Lewis et al. (2020) introduced the foundational RAG framework, combining dense retrieval and generative modeling in an integrated system. Queries are encoded via a pre-trained dense bi-encoder, relevant passages are retrieved from an external corpus, and these passages condition a sequence-to-sequence generator. RAG demonstrated substantial reductions in hallucinations and significant improvements on open-domain QA benchmarks, setting the modular retriever–generator paradigm that underpins subsequent multimodal extensions.

Muludi et al. (2024) Muludi et al. (2024) adapted RAG for document QA by segmenting long-form documents into passages indexed with dense text embeddings. Their dual-stage retrieval uses a dense retriever to surface candidates, followed by cross-encoder re-scoring for fine-grained relevance. Conditioned on top-k retrieved segments, their pipeline shows over 15% gains in exact-match accuracy on DocQA benchmarks and reduces contradictory answers.

Adjali et al. (2024) Adjali et al. (2024) proposed a Multi-Level RAG framework for Visual QA, operating hierarchically at both segment and page levels. A lightweight text retriever first narrows candidate segments, which are then re-ranked by vision-language late-interaction models fusing visual and textual features. This two-stage retrieval reduces latency by approximately 40% while improving precision on queries involving diagrams and tables.

Long et al. (2025) Long et al. (2025) introduced ReAuSE: Retrieval-Augmented Visual QA, which enriches page embeddings with semantic role labels and entity annotations. Document pages are embedded via a VLM and augmented by a lightweight NLP pipeline that tags semantic roles. An adaptive re-ranking stage refines candidate pages based on cross-modal coherence scores, reducing hallucinations and improving fidelity by 18% on complex table reasoning tasks. ReAuSE dynamically adjusts retrieval depth according to query complexity, demonstrating the impact of semantic augmentation in cross-modal alignment.

PDF WuKong (2023) PDF WuKong (2023) represents PDFs as hierarchical multimodal trees, using a tree-based encoder to capture layout hierarchy while integrating CLIP-based global embeddings and OCR-derived local tokens. This architecture achieves robustness to diverse document layouts without explicit retrieval modules, combining coarse visual context with fine textual details.

MagicLens (ICCV 2023) MagicLens (ICCV 2023) introduces a semantic lens-based retrieval strategy, projecting queries into multiple latent subspaces—‘lenses’—each specializing in tables, figures, or paragraphs. Candidate regions are retrieved via specialized vision-language encoders and aggregated for final ranking, balancing semantic specialization with retrieval breadth.

VisDoMRAG (ECCV 2024) VisDoMRAG (ECCV 2024) merges dense vision retrieval with dynamic text overlay: pages are embedded by a VLM for visual similarity, then OCR-extracted text regions are overlaid during scoring. This hybrid overlay mechanism enhances alignment of visual and textual signals, yielding performance gains on benchmarks with complex diagrams.

Comparative Analysis

Retrieval strategy VisRAG and VDocRAG rely on dense image-based retrieval to maximize visual and layout features at the expense of computational cost. RAG sets the basis for modular retrieval–generation. M3DocRAG’s late-interaction multi-modal retrieval aligns text and image tokens for fine-grained matching. Muludi et al. show that text-only dense retrieval scales efficiently on long documents. Adjali et al.’s hierarchical RAG combines text-first retrieval with vision-language re-ranking. Long et al.’s ReAuSE adds semantic role-based re-ranking. PDF WuKong avoids explicit retrieval by tree lookups, MagicLens uses parallel semantic lenses, and VisDoMRAG overlays text on vision embeddings. The author’s system uses hybrid CLIP+Col retrieval, precomputing CLIP embeddings for speed and invoking Col models for precision.

Handling visual information VisRAG, M3DocRAG, and VDocRAG process full-page images, preserving layout and graphic elements. RAG itself is text-focused but modular. DocPrompt mediates via OCR and layout tokens. Muludi et al. forgo vision entirely. Adjali et al. and Long et al. integrate vision-language interaction in re-ranking. PDF WuKong blends tree-based layout encoding with vision and OCR tokens. MagicLens retrieves through semantic lenses, and VisDoMRAG overlays text dynamically. The author’s pipeline remains fully vision-centric using CLIP and Col embeddings end-to-end.

Scalability and efficiency Dense image-based retrieval (VisRAG/VDocRAG) and multi-modal indexing (M3DocRAG) require high compute and storage. DocPrompt shifts cost to pre-training. Text-based RAG (Muludi et al.) scales via precomputed embeddings. Hierarchical retrieval (Adjali et al.) and adaptive depth (Long et al.) optimize runtime. PDF WuKong’s tree encoder reduces lookup cost. MagicLens parallel lenses improve throughput. VisDoMRAG’s overlay adds OCR overhead but focuses computation on key regions. The author’s hybrid CLIP+Col design balances speed and precision.

Generalization VisRAG and VDocRAG claim open-domain transfer from large vision-text pre-training. M3DocRAG’s modular indices support closed/open domains. DocPrompt’s layout-aware pre-training offers cross-domain robustness. Muludi et al. cover varied text

corpora. Adjali et al. and MagicLens adapt to mixed layouts. Long et al.’s ReAuSE leverages semantic roles for feature alignment. PDF WuKong’s tree representation is layout-agnostic. VisDoMRAG transfers to diagram-heavy documents. The author’s unified schema and Qwen2.5 generator ensure cross-domain QA.

Architectural modularity All RAG-based systems—Lewis et al.’s RAG, VisRAG, M3DocRAG, VDocRAG, Muludi et al., Adjali et al., Long et al.—separate retrieval and generation modules. DocPrompt is monolithic, while PDF WuKong, MagicLens, and VisDoMRAG modularize retrieval via tree encoders, semantic lenses, and overlay mechanisms, respectively. The author’s system combines CLIP and Col retrievers with a modular Qwen2.5 generator for flexible hybrid strategies.

Explainability RAG frameworks across works ground answers in retrieved evidence for traceability. VisRAG and VDocRAG include visual citations; M3DocRAG’s late-interaction highlights token alignments; Muludi et al. and Adjali et al. enable passage-level inspection; Long et al.’s semantic metrics offer interpretability; PDF WuKong’s tree view maps answers through layout hierarchy; MagicLens attributes contributions per lens; VisDoMRAG’s overlay shows text-image interplay. DocPrompt lacks explicit grounding. The author’s system retains full explainability by passing retrieved CLIP/Col embeddings with citations into Qwen2.5.

<i>System/Model</i>	<i>Retrieval Strategy</i>	<i>Handling Visual Information</i>	<i>Scalability & Efficiency</i>	<i>Generalization</i>	<i>Architectural Modularity</i>	<i>Explainability</i>
<i>VisRAG</i>	Dense, fully image-based retrieval.	Processes full pages as images to preserve layout.	High computational cost.	Generalizes from large vision-text pre-training.	Modular (retriever + generator).	Provides visual citations.
<i>VDocRAG</i>	Dense, image-based retrieval with vision-augmented pre-training.	Treats pages as images; charts & tables are vision-augmented pre-training.	Requires compute-intensive pre-training and dense indexing.	Strong generalization to new layouts.	Modular (retriever + generator).	Provides visual citations.
<i>M3DocRAG</i>	Multi-modal (image + text)	Jointly encodes images and text	Requires Approximate Nearest Neighbor	Flexibly supports open and	Modular and decoupled.	Highlights token alignments.

	retrieval with late-interaction.	to preserve visual and textual context.	(ANN) indexing to scale.	closed domains.		
DocPromp t	No inference-time retrieval; relies on continued pre-training.	Mediates information via OCR and layout tokens.	Shifts cost to the large-scale pre-training stage.	Robust cross-domain performance due to layout awareness.	Monolithic.	Lacks explicit grounding.
Lewis et al. (RAG)	Foundational RAG: dense text retrieval + generation is modular.	Primarily text-focused; architecture is modular.	Scalable and modular.	Set the paradigm for open-domain QA.	Modular (retriever + generator).	Grounds answers in retrieved passages.
Muludi et al.	Dual-stage dense text retrieval (retrieve entirely + re-rank).	Forgoes visual information on entirely (text-only).	Efficient and scalable with precomputed text embedding s.	Covers varied text corpora.	Modular.	Allows passage-level inspection.
Adjali et al.	Hierarchical multi-level retrieval (text-first, then VL re-rank).	Fuses vision and language features in the re-ranking stage.	Reduces latency via two-stage retrieval.	Adapts to mixed layouts (text and diagrams).	Modular.	Allows passage-level inspection.
Long et al. (ReAuSE)	Adds semantic role-based re-ranking.	Enriches page embeddings with semantic linguistic labels.	Optimizes runtime with adaptive retrieval depth.	Uses semantic roles for better cross-modal alignment.	Modular.	Offers interpretability via semantic metrics.
PDF WuKong	No explicit layout retrieval; relies on tree-based lookup.	Integrates layout hierarchy with CLIP and OCR tokens.	Low lookup cost due to the tree encoder.	Layout-agnostic.	Modular via the tree encoder.	Maps answers through the layout hierarchy.

MagicLen	Retrieval via semantic 'lenses' (tables, figures, text). Adapts to mixed layouts.	Uses semantic encoders for each type of visual content.	Improves throughput via parallel lenses.	Modular via lenses.	Attributes contributions per lens.
VisDoMR AG	Dense vision retrieval with dynamic text overlay.	Fuses visual and textual signals by overlaying OCR text at scoring.	Adds OCR overhead but focuses computation on key regions.	Transfers well to diagram-heavy documents.	Shows text-image interplay.
Author's System	Hybrid: CLIP (speed) + Col (precision).	Fully vision-centric using CLIP and Col embeddings.	Balances speed and precision with hybrid design.	Cross-domain QA with unified schema & Qwen2.5 generator.	Modular (retriever + generator). Full explainability via citations in generator.

Table 9 papers advantages.

In summary

the author's system synthesizes strengths from recent work. Like VisRAG and VDocRAG, it processes entire page images without OCR, preserving layout informationarxiv.orgopenaccess.thecvf.com. Like M3DocRAG, it leverages advanced multi-modal retrievers (ColQwen/ColPali) to capture both textual and visual relevance. It contrasts with DocPrompt by emphasizing an end-to-end RAG pipeline rather than relying solely on pre-trained generative modelsarxiv.org. By combining dense and late-interaction retrieval, a unified multi-dataset schema, and a strong VLM generator (Qwen2.5), the author's design aims for both flexibility and robustness, enabling scalable, explainable document QA with rich visual reasoning..

2.4 Summary and Research Gap

This review has established that while Vision-Language Models offer a powerful foundation for understanding document content, they cannot handle long-form, multi-page documents at scale. The Retrieval-Augmented Generation (RAG) framework directly addresses this limitation by combining an external knowledge source with a generator model. Our analysis of information retrieval techniques revealed that modern late-interaction neural architectures

offer the best trade-off between the efficiency of dense retrieval and the effectiveness of cross-encoders.

However, a significant research gap remains: most RAG systems are text-based and thus fail to leverage the rich visual and structural information present in documents. This leads to a critical trade-off where systems are either scalable but visually unaware, or visually aware but unscalable.

Therefore, this research is positioned to address this gap by proposing a hybrid framework that integrates a **visually-aware, page-level retrieval mechanism** with a **high-performance VLM generator**. The following chapter details the architecture of this proposed solution.

Chapter 3: The Proposed Solution

3.1 Solution Methodology

After thoroughly analyzing the research problem and identifying the limitations of existing DocVQA systems through an extensive literature review, the next step was to design and implement a concrete solution that directly addresses these challenges.

This section presents the methodology followed in developing the proposed system. The approach integrates concepts from retrieval-augmented generation (RAG) with the capabilities of vision-language models (VLMs), aiming to achieve both long-context reasoning and layout awareness. The system was implemented in a modular fashion to allow for flexibility, scalability, and comparative evaluation.

The methodology is guided by the following principles:

- Leverage existing pre-trained models to reduce training cost.
- Introduce layout-sensitive retrieval to enhance page selection.
- Evaluate both few-shot and fine-tuned variants for comprehensive assessment.

The following subsections outline the detailed system architecture, the reasoning behind each design choice, and how the proposed solution differs from or improves upon existing approaches.

3.1.1 Problem Scope Refinement

The problem of DocVQA is inherently multimodal and complex due to:

- The **length** and **multi-page** nature of documents.
- The importance of **layout and visual context** (e.g., tables, headers, spatial clues).
- The limited **input size** of most Vision-Language Models (VLMs).

Thus, the proposed solution narrows the focus to:

- Retrieval-augmented approaches (rather than monolithic VLMs).
- Page-level granularity instead of full-document encoding.
- End-to-end trainability with optional fine-tuning.

This refined scope directly informs the system architecture, described next

3.1.2 System Architecture Overview (M3DocRAG Framework)

The proposed solution is built around **M3DOCRA**G (Multi-modal, Multi-page, Multi-Document Retrieval-Augmented Generation), a novel framework designed to handle diverse document contexts (closed/open-domain), question types (single/multi-hop), and evidence modalities (text, charts, figures). As illustrated in Figure 18, M3DOCRA operates in three stages: (1) document embedding, (2) page retrieval, and (3) question answering. Unlike prior DocVQA datasets (Figure 19), which focus on single-document queries, M3DOCRA scales to open-domain settings across thousands of documents..




Figure 13 M3DOCRA framework retrieves relevant

In contrast to previous DocVQA datasets that have questions that are specific to a single provided PDF (e.g., “What was the gross profit in the year 2009?”), M3DOCVQA has information-seeking questions that benchmark open-domain question answering capabilities across more than 3,000 PDF documents (i.e., 40,000+ pages)




Figure 14 Comparison of existing DocVQA datasets and M3DOCVQA

M3DOCRA operates in three stages:

1. **Document Embedding:**
we convert all document pages into RGB images and extract visual embeddings (e.g., via ColPali) from the page images.
2. **Page Retrieval:**
we retrieve the top-K pages of high similarity with text queries (e.g., MaxSim operator for ColPali). For the open domain setting, we create approximate page indices, such as inverted file index (**IVF**), for faster search
3. **Question Answering:** we conduct visual question answering with MLM to obtain the final answer. Please also see Fig.20 for the detailed illustration of the framework. M3DOCRA can

flexibly handle DocVQA in both closed domain (i.e., a single document) and open-domain (i.e., a large corpus of documents) settings.




Figure 15 M3DOCRAG framework consists of three stages: document embedding page retrieval

3.1.3 Problem definition

We define a corpus of documents as

$$C = \{D_1, D_2, \dots, D_M\}$$

Equation 1 corpus of documents

where M is the total number of documents, and each document D_i consists of a set of pages, p_i represented as RGB images. From the documents in C we construct a global set of page images:

$$P = \bigcup_{i=1}^M P_i = \{p_1, p_2, \dots, p_N\},$$

Equation 2 global set of page images

where each p_j represents an individual page image, and N is the total number of page images across all documents in C

$$\left(i.e., N = \sum_{i=1}^M |P_i| \right)$$

Equation 3 number of page

The objective of M3DOCRAG is to accurately answer a given question q using the multimodal information available in the corpus of documents C .

First, we identify P_q^K , the top K ($\ll N$) pages that are most relevant to answering the query q from the global page set P . Then, we obtain the final answer using a question answering model that takes the retrieved page images P_q^K and the query q as inputs.

The problem of question answering can be categorized into two settings with different document context sizes:

- **Closed-domain question answering** – The query q should be answerable from a given single document D_i . The retrieval model outputs the top K relevant page images P_q^K , from the page images P_q^K of the document D_i
- **Open-domain question answering** – The query q may require information from single or multiple documents within the entire document corpus C . The retrieval model outputs the top K relevant page images P_q^K from the entire set of page images P .

3.1.4 Document Embedding

In **M3DOCRAg**, both the textual query q and page images P are projected into a shared multi-modal embedding space using **ColPali**. ColPali is a multi-modal retrieval model based on a **late interaction mechanism**, which encodes the text and image inputs into unified vector representations and retrieves the top K most relevant images. ColPali adopts both the training objective and similarity scoring from **ColBERT**, which utilizes a shared architecture to encode either textual or visual inputs.

In this framework, each page $p \subseteq P_k$ of a document D_k is treated as a single image with fixed dimensions (width \times height).

From an image of a page, we extract a dense visual embedding:

$$E_p \in \mathbb{R}^{n_v \times d},$$

where n_v represents the number of visual tokens per page (which remains constant across all pages), and d denotes the embedding dimension (e.g., 128). For a textual query q , we similarly obtain an embedding:

$$E_q \in \mathbb{R}^{n_q \times d},$$

where n_q is the number of text tokens.

For efficiency, we treat each page of a document independently. This allows us to flatten all pages in the document corpus C into a single page-level embedding tensor:

$$E_C \in \mathbb{R}^{N \times n_v \times d},$$

where N represents the total number of pages in the entire document corpus, n_v is the number of visual tokens per page, and d is the embedding dimension.

M3DOCRAg can flexibly adapt to different retrieval settings, such as:

- A single-page document ($N = 1$)
- A single document with multiple pages (e.g., $N = 100$)
- A large corpus of multi-page documents (e.g., $N > 1,000$)

3.15 Page Retrieval

The relevance between the query q and a page p is computed using the **MaxSim** score $s(q, p)$:

$$s(q, p) = \sum_{i=1}^{n_q} \max_{j \in [n_p]} \langle E_q^{(i,\cdot)}, E_p^{(j,\cdot)} \rangle$$

Equation 4 MaxSim score

where (\cdot) denotes the dot product, and $(E_{i,\cdot} \in R^d)$ denotes the (i) -th row (vector) of the embedding matrix $(E \in R^{n \times d})$. We then identify (P_K^q) , the top (K) ($\ll N$) pages that are most relevant to answering the query (q) ; i.e., we search (K) pages scoring highest $(s(q, p))$. That is,

$$P_q^K = \{p_q^1, p_q^2, \dots, p_q^K\} = \arg \max_{\text{top-}k, p \in P} s(q, p)$$

Equation 5 Top-K page selection

Approximate indexing for open-domain page retrieval

Searching over a large document corpus can be time-consuming and computationally expensive. For faster retrieval, we build offline page indices using approximate nearest neighbor (ANN) search based on **Faiss**.

We use exact search for closed-domain retrieval, and for open-domain retrieval, we apply the **inverted file index (IVF)** (IVFFlat in Faiss), which can reduce page retrieval time from 20 seconds per query to less than 2 seconds when searching across 40K pages. (See Section 5.3 for a detailed speed–accuracy trade-off comparison.)

3.1.6 Question Answering

We run visual question answering by giving the text query (q) and retrieved page images (P_K^q) to a multi-modal language model to obtain the final answer. For this, we employ multi-modal language models (e.g. Qwen2-VL) that consist of a visual encoder (Enc^{Vis}) and a language model LM. The visual encoder takes K-retrieved page (P_K^q) as inputs and outputs visual embeddings (different from ColPali encoder’s outputs). The language model takes the visual embeddings and text embeddings of query (q) as inputs and outputs the final answer (a) in the autoregressive manner:

$$a = \text{LM}(\text{Enc}_{\text{Vis}}(P_q^K), q)$$

Equation 6 Question answering output

The selection of models and tools was guided by empirical research and the need to balance performance, efficiency, and scalability. Below, we detail the rationale behind each critical component




Figure 13 TextRAG (left) vs. VisRAG (right)

3.1.7 Models and Tools Selection

Following the architectural design of the proposed Document Visual Question Answering (DocVQA) system, a key decision point emerges: selecting the models and tools that form the core of the pipeline. This stage is fundamental to the performance, scalability, and practical impact of the entire solution. The selection process was neither arbitrary nor limited to replicating previous work; rather, it was grounded in a comprehensive analysis of the strengths and limitations of available methods, especially under the lens of retrieval-augmented generation (RAG) applied to multimodal document understanding.

The primary goal was to strike a balance between:

- High performance in multi-modal understanding and retrieval
- Processing efficiency and speed
- Flexibility to operate under both open-domain and closed-domain settings
- Scalability and extensibility for future enhancements

The following sections provide a detailed explanation of **why each model or tool was selected**, along with a discussion of how each component contributes to the overall objectives of the proposed system.

Retrieval Model Selection

In retrieval-augmented DocVQA systems, the retrieval module plays a decisive role: it determines which pages from a potentially massive document collection are forwarded to the reasoning stage. A weak retriever propagates irrelevant context to the reader model, severely degrading final performance. Conversely, an accurate and efficient retriever directly enhances answer relevance, system latency, and scalability.

Thus, the chosen retrieval strategy must strike a delicate balance across the following criteria:

- **Layout-aware understanding** to handle complex document structures like tables and figures.
- **Scalability to open-domain settings**, where retrieval spans tens of thousands of pages.
- **Inference efficiency** to keep query-time latency acceptable (<2s/query).
- **Multimodal support**, integrating both text and visual features into the retrieval process

Comparison of Retrieval Paradigms

To inform our selection, we conducted a structured comparison of four main retrieval paradigms using state-of-the-art models:

Paradigm (Example)	Interaction Pattern	Evidence Modality	Pros	Cons
Sparse lexical retrieval (BM25)	Keyword overlap scored offline	OCR text only	Fast index construction; no GPU cost	Ignores layout and images; weak on synonymy
Dense bi- encoder (DPR, CLIP dual encoder)	<i>Early interaction:</i> query and page each → one vector; similarity = dot product	Text or image	Pre-compute page vectors; sub- second ANN search	One vector often loses fine- grained token or region signals
Cross-encoder reranker (BERT, BLIP-2 Q-I cross-attn)	<i>Full interaction:</i> query + page concatenated in one transformer	Text + image	State-of-the-art ranking accuracy	Cannot be pre- indexed; linear in corpus size— impractical for first-stage retrieval

Late-interaction <i>(ColBERT, ColPaLi family)</i>	Page tokens pre-encoded; query tokens encoded at run-time; similarity = MaxSim over token pairs	Text (ColBERT) or Text + image (ColPaLi)	Token-level matching improves recall; pages still pre-indexed; good speed/quality balance	Larger index size than bi-encoder; extra MaxSim cost
---	---	--	---	--

Table 10 retrieval comparison

Key observations

- **Sparse keyword** methods are efficient but blind to layout and visual semantics—an unacceptable limitation for DocVQA.
- **Cross-encoders** deliver the best ranking quality yet are prohibitively slow for corpora exceeding a few hundred pages.
- **Dense bi-encoders** scale well but compress each page into a single vector, discarding many region-level cues required for table or figure questions.
- **Late-interaction architectures** preserve token-level (or patch-level) detail while allowing offline indexing, offering a promising middle ground.

Selection criteria for this study

1. **Precision on layout-sensitive queries** (tables, figures, multi-column text).
2. **Retrieval latency below two seconds for 40 000 pages** (open-domain setting).
3. **Compatibility with a multimodal encoder** so that both visual and textual evidence are indexed jointly.

Embedding Methods Comparison: LAION CLIP, ColPaLi, and ColQwen

To deepen our retrieval strategy decision, we compared embedding approaches of three leading models.

1.LAION CLIP (Contrastive Language-Image Pre-Training)

Embedding Approach

- Uses contrastive learning to align text and image in a shared space
- Produces a single dense vector per input
- Pretrained on LAION-5B for broad semantic coverage

Process

- Text: Encoded by a Transformer
- Image: Encoded by a Vision Transformer (e.g., ViT-B/32)
- Retrieval: Cosine similarity between embeddings

Characteristics

- **Strengths:** Simple, effective semantic matching
- **Limitations:** Not layout-aware; OCR required for documents
- **Best for:** Image-text matching in simple settings

2. ColPaLi (Contextualized Late Interaction over PaliGemma)

Embedding Approach

- Uses PaliGemma-3B with late-interaction
- Generates ~1,030 vectors per page (patches + instruction tokens), each 128D

Process

- Image split into 1,024 patches; processed by SigLIP + Gemma-2B
- Query tokens embedded into same 128D space
- Retrieval via **MaxSim** scoring between tokens and patches

Characteristics

- **Strengths:** Fine-grained matching; layout-sensitive; efficient retrieval (~0.39s/page)
- **Limitations:** Higher index size; restrictive license
- **Best for:** Complex, visual-rich documents

3. ColQwen (ColQwen2 with Qwen2-VL)

Embedding Approach

- Builds on ColPaLi with Qwen2-VL-2B
- Produces ~700–768 vectors per page, each 128D (efficiency boost)

Process

- Vision + language embeddings created per patch
- Retrieval using MaxSim, like ColPaLi

Characteristics

- **Strengths:** Efficient; Apache 2.0 license; better resolution handling
- **Limitations:** Slightly coarser than ColPaLi
- **Best for:** Scalable, production-ready systems

FEATURE	LAION CLIP	COLPALI	COLQWEN
EMBEDDING TYPE	Single-vector	Multi-vector	Multi-vector
VECTORS PER PAGE	1	~1,030	~700–768
VECTOR DIMENSIONALITY	512/768	128	128
INPUT MODALITY	Text or image	Page images	Page images
STRENGTHS	Semantic generality	Fine-grained, layout-aware	Efficient, scalable
WEAKNESSES	Needs OCR	Larger storage, license	Slightly less precise
LICENSE	Open-source	Restrictive	Apache 2.0

Justifying the Use of ColPaLi as the Baseline

Among the surveyed retrieval architectures, **ColPaLi** stood out as the most balanced and robust solution for our task. The decision to adopt it as the baseline retriever in our system was driven by both **quantitative evidence** and **architectural advantages**, rather than merely replicating the original M3DocRAG pipeline.




Figure 14 ColPaLi simplifies document retrieval

A. Architectural Strengths of ColPaLi

- **Multimodal Late Interaction:** Unlike dense models, ColPaLi preserves fine-grained information across both text and image modalities through a patch-level MaxSim scoring mechanism.
- **Layout-Aware Retrieval:** Token-to-patch alignment allows ColPaLi to handle visually structured inputs like tables and diagrams with higher fidelity.
- **Efficient Indexing at Scale:** With Faiss IVF-based indexing and token pooling, ColPaLi achieves sub-2s retrieval latency over 40K pages without sacrificing retrieval quality.

B. Empirical Performance (ViDoRe Benchmark)

To avoid anecdotal justification, we anchored our model choice in empirical findings from the **ViDoRe benchmark**—a comprehensive retrieval evaluation framework spanning.

- **10 Datasets:** Covering financial documents (TAT-DQA), scientific figures (ArxivQA), French tables (TabFQuAD), and more.
- **3 Metrics:** Focusing primarily on **nDCG@5** for retrieval quality, with a strict **latency constraint (<2s/query)**.
- **4 Paradigms:** Comparing **sparse**, **dense**, **cross-encoder**, and **late-interaction** models.

MODEL	TABFQUAD (TABLES)	SHIFT PROJECT (FRENCH DOCS)	OVERALL (NDCG@5)
CLIP (DENSE)	58.1	62.3	64.5
SIGLIP (DENSE)	61.2	65.0	66.8
BLIP-2 (CROSS)	74.0	70.1	71.5
COLPALI	83.9	73.2	77.8

Table 11 These results of the evaluation

These results confirm that ColPaLi not only outperforms dense baselines in layout-sensitive queries but also maintains real-time inference capabilities, capable of handling structured, multilingual, and visually rich documents. making it the most suitable default retriever.

Key Insight: No existing solution performed optimally across all axes until **ColPaLi**.

Quantitative Comparison

Retrieval Performance on ViDoRe (nDCG@5):

	ArxivQ	DocQ	InfoQ	TabF	TATQ	Shift	AI	Energy	Gov.	Health.	Avg.
Unstructured text-only											
- BM25	-	34.1	-	-	44.0	59.6	90.4	78.3	78.8	82.6	-
- BGE-M3	-	28.4 _{↓5.7}	-	-	36.1 _{↓7.9}	68.5 _{↑8.9}	88.4 _{↓2.0}	76.8 _{↓1.5}	77.7 _{↓1.1}	84.6 _{↑2.0}	-
Unstructured + OCR											
- BM25	31.6	36.8	62.9	46.5	62.7	64.3	92.8	85.9	83.9	87.2	65.5
- BGE-M3	31.4 _{↓0.2}	25.7 _{↓11.1}	60.1 _{↓2.8}	70.8 _{↑24.3}	50.5 _{↓12.2}	73.2 _{↑8.9}	90.2 _{↓2.6}	83.6 _{↓2.3}	84.9 _{↑1.0}	91.1 _{↑3.9}	66.1 _{↑0.6}
Unstructured + Captioning											
- BM25	40.1	38.4	70.0	35.4	61.5	60.9	88.0	84.7	82.7	89.2	65.1
- BGE-M3	35.7 _{↓4.4}	32.9 _{↓5.4}	71.9 _{↑1.9}	69.1 _{↑33.7}	43.8 _{↓17.7}	73.1 _{↑12.2}	88.8 _{↑0.8}	83.3 _{↓1.4}	80.4 _{↓2.3}	91.3 _{↑2.1}	67.0 _{↑1.9}
Contrastive VLMs											
Jina-CLIP	25.4	11.9	35.5	20.2	3.3	3.8	15.2	19.7	21.4	20.8	17.7
Nomic-vision	17.1	10.7	30.1	16.3	2.7	1.1	12.9	10.9	11.4	15.7	12.9
SigLIP (Vanilla)	43.2	30.3	64.1	58.1	26.2	18.7	62.5	65.7	66.1	79.1	51.4
Ours											
SigLIP (Vanilla)	43.2	30.3	64.1	58.1	26.2	18.7	62.5	65.7	66.1	79.1	51.4
BiSigLIP (+fine-tuning)	58.5 _{↑15.3}	32.9 _{↑2.6}	70.5 _{↑6.4}	62.7 _{↑4.6}	30.5 _{↑4.3}	26.5 _{↑7.8}	74.3 _{↑11.8}	73.7 _{↑8.0}	74.2 _{↑8.1}	82.3 _{↑3.2}	58.6 _{↑7.2}
BiPali (+LLM)	56.5 _{↓2.0}	30.0 _{↓2.9}	67.4 _{↓3.1}	76.9 _{↑14.2}	33.4 _{↑2.9}	43.7 _{↑17.2}	71.2 _{↓3.1}	61.9 _{↓11.7}	73.8 _{↓0.4}	73.6 _{↓8.8}	58.8 _{↑0.2}
ColPali (+Late Inter.)	79.1 _{↑22.6}	54.4 _{↑24.5}	81.8 _{↑14.4}	83.9 _{↑7.0}	65.8 _{↑32.4}	73.2 _{↑29.5}	96.2 _{↑25.0}	91.0 _{↑29.1}	92.7 _{↑18.9}	94.4 _{↑20.8}	81.3 _{↑22.5}

Table 12 Comprehensive evaluation of baseline models

C. Multimodal Fusion & Layout Awareness

- Uses **PaliGemma-3B** for projecting both text and visual elements (e.g., tables, charts) into a unified semantic space.
- Preserves **token-to-patch alignment**, essential for layout-sensitive QA tasks.

D. Scalability & Speed

- Integrates with **Faiss IVF indexing**, reducing retrieval time from 20s to under 2s on corpora with 40K+ documents.
- Optimized via **token pooling** and **dimensionality reduction**, decreasing memory footprint significantly without performance loss.

Design Rationale

Rather than relying on ColPaLi just because it was used in M3DocRAG, our choice is:

- **Evidence-driven**, based on multi-domain benchmarks.
- **Performance-based**, as it achieved state-of-the-art results with practical latency.
- **Architecture-aligned**, offering late interaction and multimodal capabilities essential for DocVQA.

Comparative Experiments in Chapter 4

To validate this choice further, Chapter 4 will present a comparative evaluation of:

- **ColPaLi (Late Interaction, Multimodal)**
- **ColQwen (Multimodal Late Interaction)**
- **CLIP & SigLIP (Dense Bi-Encoder Baselines)**

All models will be tested under **identical ANN search settings** using a different dataset than M3DocRAG (our own evaluation set), measuring:

- Retrieval accuracy (nDCG@5, recall)
- Latency
- Resource consumption

Conclusion: Why We Chose ColPaLi

Building on both our analysis and empirical evidence from the original **M3DocRAG paper**, we selected **ColPaLi** as the primary retriever for our system. This choice is not arbitrary but grounded in **benchmark results** and architectural strengths that align with the unique challenges of multimodal document retrieval.

ColPaLi demonstrates the following advantages:

- **Late-interaction design**, enabling fine-grained token-patch alignment and preserving layout information.
- **Multimodal fusion capability**, thanks to PaliGemma's encoder that effectively integrates text and visual signals into a unified embedding space.
- **High scalability and retrieval efficiency**, supported by FAISS-based indexing and embedding compression (e.g., IVF+FlatIP).

These benefits are empirically validated in **Table 10 of the M3DocRAG paper**, where ColPaLi v1 outperforms ColQwen v0.1 on the **M3DocVQA** benchmark with a substantial margin ($F1 = 36.5$ vs. 32.1), highlighting its strength in layout-sensitive, real-world documents. While ColQwen performs slightly better on **MP-DocVQA** and **MMLongBench-Doc**, the authors themselves select **ColPaLi as their default retrieval model**, citing its **robust performance in the main benchmark (M3DocVQA)** and architectural advantages derived from PaliGemma.

Ret. Models	M3DocVQA	MMLongBench-Doc	MP-DocVQA
	F1 (\uparrow)	Acc (\uparrow)	ANLS (\uparrow)
ColPali v1	36.5	21.0	0.84
ColQwen v0.1	32.1	21.5	0.86

Table 13 Comparison between ColPaLi and ColQwen

“We find that ColPaLi achieves significantly better performance in M3DocVQA... Thus, we use ColPaLi as our default retrieval model.” — M3DocRAG, Sec. 4

In addition, as part of our own comparative study in **Chapter 4**, we benchmark the following retrieval models under identical ANN settings:

- Two late-interaction retrievers: **ColPaLi** and **COLQwen**
- Two dense bi-encoders: **CLIP** and **SigLIP**

This will allow us to empirically validate the original authors' claims and further assess the generalization of ColPaLi's performance to a different dataset. Until then, we adopt ColPaLi as the baseline for its demonstrated balance between **retrieval quality, speed, and multimodal understanding**.

ColQwen Model (Illuin Technology)

ColQwen is a multimodal late-interaction retriever built on the Qwen2-VL vision-language backbone. It extends Qwen2-VL-2B by generating ColBERT-style multi-vector embeddings for document images and text queries [40]. In practice, a Qwen2-VL Vision Encoder splits each document page image into up to 768 patches (un-resized, dynamic resolution), encodes them with a ViT-like transformer, and feeds the patch embeddings into the Qwen2 language model. The Qwen2 Language Model (2B) also serves as the query encoder, producing contextualized token embeddings for the text query. During retrieval, a late-interaction (“MaxSim”) is computed between each query token and all image patch vectors, summing the top similarities [41]. This yields a rich matching score akin to ColBERT. ColQwen's outputs are 128-dimensional vectors per token/patch [40].

Figure: ColQwen's architecture. A Qwen2-VL vision encoder produces patch embeddings from the document image, which are fed into a Qwen2 language model (purple) to yield multi-vector image representations. The Qwen2 text model (green) encodes the query. A ColBERT-style late interaction then matches query tokens to document patch embeddings [41].

In summary, ColQwen's architecture mirrors ColPali but uses the Qwen2-VL backbone. It has ~2 billion parameters and projects outputs to 128-dimensional space. Compared to ColPali's PaliGemma-3B backbone, ColQwen's Qwen2-VL is ~1B parameters smaller and uses a smaller patch grid (768 vs 1024), reducing memory and compute [40].

Training & Data

The model is trained with a contrastive retrieval objective on image–text (query–document)

pairs. Training uses a curated document dataset: ~127K query–page pairs (63% real academic pairs, 37% synthetic using VLM-generated questions) [42]. Images are fed at high resolution (up to 768 patches), and the final projection layer to 128-D is randomly initialized. During fine-tuning, LoRA adapters (rank=32) are applied to all Qwen2-VL transformer layers, and an AdamW 8-bit optimizer (paged AdamW) is used with bfloat16 precision [43]. Models are trained for one epoch (batch size 32) on 8 GPUs with linear LR decay and 2.5% warmup. The training data is English-only (though the Qwen2 LM itself was pretrained on multilingual text). Notably, ColQwen’s code (via the colpali-engine library) is fully open-source under MIT/Apache licenses [44].

Performance (Benchmarks)

ColQwen targets visually-rich document retrieval (e.g. PDF pages), not general vision tasks. In the ViDoRe benchmark (Visual Document Retrieval), ColQwen2 outperforms the earlier ColPali model by several points: for example, it yields about +5 NDCG@5 improvement [45]. This means more accurate retrieval of relevant pages given a query. In practice, ColQwen achieves higher retrieval accuracy on tasks involving tables, figures and complex layouts, while greatly simplifying ingestion (no OCR or layout parsing is needed).

Model Cards, Code and Papers

- **Paper:** *ColPali: Efficient Document Retrieval with Vision Language Models* [46].
 - **Model Card:** HuggingFace “vidore/colqwen2-v0.1” [40].
 - **Code:** [44]
 - **Datasets:** ViDoRe Benchmark [45].
 - **Licensing:** Qwen2-VL (Apache-2.0); ColQwen code (MIT) [44].
-

LAION CLIP Models

LAION’s CLIP models are open-source contrastively trained vision-language encoders. They follow the original CLIP architecture: a text encoder (Transformer) and an image encoder (ViT or ResNet) trained to align image–text pairs [47]. For example, the CLIP-ViT-L/14 and ViT-H/14 models (from laion/CLIP-ViT-* on HuggingFace) use ViT backbones for images and a standard 12-layer Transformer for text [48]. These models are trained on the English subset of LAION-5B: 2 billion image–caption pairs scraped from the web. Training is done with OpenCLIP code on large GPU clusters (JUWELS and stability.ai clusters), seeing tens of billions of examples in total [49].

Figure: CLIP’s contrastive training: image and text encoders (blue blocks) jointly embed data, and a contrastive loss pulls matching pairs together [50].

LAION CLIP uses a batch-size-contrastive objective: each minibatch of image–text pairs is fed through separate encoders, and a symmetric cross-entropy (InfoNCE) loss maximizes cosine similarity of true pairs and minimizes it for mismatched pairs [47]. The LAION-2B training data is unfiltered (uncurated), requiring caution (potential NSFW content) [48].

Benchmarks & Performance

Notably, the ViT-H/14 version achieves ~78.0% top-1 accuracy on ImageNet (zero-shot) and 73.4% Recall@5 on MS-COCO image-text retrieval. The ViT-L/14 model also performs well

(\approx 75.3% ImageNet, 71.1% COCO R@5). Both models yield \sim 94% and 92.9% Recall@5 on Flickr30k image-text retrieval, respectively [49]. These results (from LAION’s blog) surpass earlier CLIP benchmarks for models of similar size, benefiting from much larger training data [47].

Training & Data

CLIP models were trained on \sim 34B image–text pairs with batch sizes up to 79K [47][49]. They use standard ViT encoders, trained in mixed precision on web-scale image-text data, without adapters or tricks.

Model Cards, Code and Papers

- **Model Cards:** HuggingFace: laion/CLIP-ViT-L-14-laion2B, laion/CLIP-ViT-H-14-laion2B [48].
- **Code:** OpenCLIP GitHub [50].
- **Data:** LAION-5B Dataset [51].
- **Paper:** Radford et al. (2021), CLIP [46]; LAION blog/paper: Beaumont et al. (2022) [47].

Comparison of ColQwen and LAION CLIP

- **Purpose & Tasks:** ColQwen specializes in document/page retrieval using late interaction; LAION CLIP is general-purpose image-text retrieval.
- **Architecture:** ColQwen uses Qwen2-VL with late interaction; CLIP uses dual unimodal encoders and global vectors [40][47].
- **Performance:** ColQwen outperforms ColPali on ViDoRe (\sim +5 NDCG@5) [45]; CLIP reaches SOTA in general image benchmarks (ImageNet, COCO, Flickr30k) [49].
- **Design:** ColQwen uses LoRA and 8-bit AdamW [43]; CLIP uses massive data and batch sizes without fine-tuning tricks [47].
- **Open-Source:** Both are MIT/Apache licensed [44][50].

ColPaLi ARCHITECTURE

Vision-Language Models. Encouraged by their strong document understanding capabilities, we propose adapting recent VLMs for retrieval. The key concept is to leverage the alignment between output embeddings of text and image tokens acquired during multi-modal fine-tuning. To this extent, we introduce ColPali, a Paligemma-3B extension that is capable of generating ColBERT-style multivector representations of text and images (Figure 1). PaliGemma-3B is a strong candidate due to its small size, the many released checkpoints fine-tuned for different image resolutions and tasks, and the promising performances on various document understanding benchmarks. We add a projection layer to map each of the language model’s output token embeddings (whether from text or image tokens) to a vector space of reduced dimension $D = 128$ as used in the ColBERT paper (Khattab & Zaharia, 2020) to keep lightweight bag-of-embedding representations.

Late Interaction. Given query q and document d , we denote as $E_q \in R^{N_q \times D}$ their respective multi-vector representation in the common embedding space R^D where N_q and N_d are respectively the number of vectors in the query and in the document page embeddings. The late interaction operator, $LI(q,d)$, is the sum over all query vectors $E_q(j)$ of its maximum dot product $\langle \cdot | \cdot \rangle$ with each of the N_d document embedding vectors $E_d(1:N_d)$

Equation 7 Late Interaction

$$LI(q, d) = \max_{j \in [1, N_d]} \langle E_q(i) | E_d(j) \rangle$$

Contrastive Loss. The Late Interaction operation is fully differentiable, enabling backpropagation. Let a batch $\{q_k, d_k\}_{k \in [|1, b|]}$ composed of b query-page pairs, where for all $k \in [|1, b|]$ the document page d_k is the document corresponding to query

q_k . Following Khattab & Zaharia (2020), we define our in-batch contrastive loss L as the softmaxed cross-entropy of the positive scores $s_k^+ = LI(q_k, d_k)$ w.r.t. to the maximal in-batch negative scores $s_k^- = \max_{l \neq k} LI(q_k, d_l)$

Equation 8 Contrastive Loss.

$$L = -\frac{1}{b} \sum_{k=1}^b \log \left(\frac{\exp(s_k^+)}{\exp(s_k^+) + \exp(s_k^-)} \right) = \frac{1}{b} \sum_{k=1}^b \log(1 + \exp(s_k^- - s_k^+)) \backslash tag{2}$$

LATENCIES & MEMORY FOOTPRINT

Online Querying. (R2) Logically, querying latencies differ between ColPali and a BGE-M3 embedding model. For BGE, encoding takes about 22 ms for 15 tokens, while encoding a query with ColPali’s language model takes about 30 ms¹². For smaller corpus sizes, computing the late interaction operation induces marginally small overheads (≈ 1 ms per 1000 pages in the corpus), and the cosine similarity computation between bi-encoder vectors is even faster. Optimized late interaction engines (Santhanam et al., 2022; Lee et al., 2023) enable to easily scale corpus sizes to millions of documents with reduced latency degradations.

Offline Indexing. (R3) Standard retrieval methods using bi-encoders represent each chunk as a single vector embedding, which is easy to store and fast to compute. However, processing a PDF to get the different chunks is the most time-consuming part (layout detection, OCR, chunking), and using captioning to handle multimodal data will only exacerbate this already lengthy process. On the other hand, ColPali directly encodes pages from their image representation. Although the model is larger than standard retrieval encoders, skipping the preprocessing allows large speedups at indexing¹³ (Figure 2). As pages are embedded end-to-end in single forward pass, the VRAM usage depends exclusively on the sequence length (number of patches per image) which is fixed as well, enabling efficient batching strategies to fully leverage hardware acceleration. ColPali also benefits from most

LLM efficiency improvements introduced in the ecosystem such as Flash Attention (Dao, 2023).

Storage Footprint. Our method requires storing a vector per image patch, along with 6 extra text tokens “Describe the image” concatenated to image patches. We project each PaliGemma vector to a lower dimensional space ($D = 128$) to maximize efficiency, leading to a memory footprint of 257.5 KB per page (subsection B.3). Importantly, the memory footprint of the naive ColBERT indexing strategy can be drastically improved through compression and clustering mechanisms (Santhanam et al., 2022; Clavi’ e et al., 2024).

Token pooling. Token pooling (Clavi’ e et al., 2024) is a CRUDE-compliant method (document addition/deletion-friendly) that aims to reduce the amount of multi-vector embeddings. For ColPali, many image patches share redundant information, e.g. white background patches. By pooling these patches together, we can reduce the amount of embeddings while retaining most information. Retrieval performance with hierarchical mean token pooling on image embeddings is shown in Figure 3 (left). With a pool factor of 3, the total number of vectors is reduced by 66.7% while 97.8% of the original performance is maintained. We note that the Shift dataset—composed of the most text-dense documents—is a clear outlier, showcasing more information dense documents contain less redundant patches and may be prone to worse performance degradation with such pooling techniques.




Figure 15 (Left: Token Pooling) Relative performance degradation when reducing the number of stored embeddings per document. (Right: Interpretability) For each term in a user query

Question Answering Model Selection:

Rationale for Selecting Qwen2.5-VL for Answer Generation

While the original **M3DocRAG paper** adopted **Qwen2-VL (7B)** as the answer generation model, our implementation utilizes a more recent and **significantly improved version—Qwen2.5-VL**. The choice to upgrade was driven by several important factors:

- **Advancements in Model Architecture:** Qwen2.5-VL introduces architectural enhancements over Qwen2-VL, including stronger vision encoders, improved multi-modal alignment, and expanded context handling. These updates are specifically tailored for **document-level understanding**, a critical requirement in DocVQA tasks.
- **Superior Performance Across Benchmarks:** According to the **Qwen2.5 Technical Report**, the model consistently outperforms both its predecessor (Qwen2-VL) and several other state-of-the-art models, including **InternVL2** and **Idefics2/3**, on a variety of multimodal reasoning tasks. This includes document-rich datasets, long-context visual reasoning, and layout-heavy scenarios, all of which are highly representative of our DocVQA domain.
- **Open Weights and Scalability:** Unlike proprietary models (e.g., Gemini Pro, GPT-4V), Qwen2.5-VL is open-weight and can be integrated efficiently into local RAG pipelines. This ensures reproducibility, customization, and full control over inference behavior.

Therefore, Qwen2.5-VL was chosen as our default answer generation module due to its **state-of-the-art performance, open-source availability, and compatibility with the large-context, visually dense nature of document QA**.

“Qwen2.5-VL demonstrates superior reasoning in long documents and structured inputs compared to InternVL2 and Idefics3.” — Qwen2.5 Technical Report (2024)

Nonetheless, as part of our evaluation strategy, we will benchmark Qwen2.5-VL against other prominent models (such as the original Qwen2-VL and other vision-language baselines) in **Chapter 4**, under identical test conditions. This comparative analysis will ensure that its

adoption is justified not only by theoretical advantages but also by empirical results specific to our task and dataset.

The model’s distinguishing characteristics are summarized below:

KEY CAPABILITY	RELEVANCE TO DOCVQA	EVIDENCE
OMNIDOCUMENT PARSING	Recognises printed text, handwriting, tables, charts, chemical formulae, music notation, and multilingual content—exactly the heterogeneity found in real-world PDFs.	Tops DocVQA-test (96.4 %) and sets new state-of-the-art on OmniDocBench (en/zh).
PRECISE OBJECT GROUNDING	Supports absolute-coordinate and JSON outputs for “where-is-it?” queries, enabling fine-grained spatial reasoning (e.g., locate a cell in a financial table).	87.1 % on ScreenSpot (grounding GUI elements) and 93.6 % on CountBench (counting objects).
ULTRA-LONG VIDEO UNDERSTANDING	Extends native dynamic resolution into the temporal domain, allowing the model to answer questions about hours-long videos embedded in reports or slide decks.	74.6 mIoU on MLVU and 50.9 mIoU on Charades-STA—even at 768 frames / 24 576 tokens.
ENHANCED AGENT FUNCTIONALITY	Combines grounding, reasoning, and decision-making to automate UI tasks on desktop or mobile—useful for future workflow agents that must open, scroll, and query documents.	Best open-weight scores on Android Control (High-EM 67.4 %, Low-EM 93.7 %) and MobileMiniWob++ (68 % SR).

Table 14 qwen2.5 advantages




Figure 17 The Qwen2.5-VL framework demonstrates the integration of a vision encoder and a language

Model Architecture

The overall model architecture of Qwen2.5-VL consists of three components: Large Language Model: The Qwen2.5-VL series adopts large language models as its foundational component. The model is initialized with pre-trained weights from the Qwen2.5 LLM. To better meet the demands of multimodal understanding, we have modified the 1D RoPE (Rotary Position Embedding) to our Multimodal Rotary Position Embedding Aligned to Absolute Time. Vision Encoder: The vision encoder of Qwen2.5-VL employs a redesigned Vision Transformer (ViT) architecture. Structurally, we incorporate 2D-RoPE and window attention to support native input resolutions while accelerating the computation of the entire visual encoder. During both training and inference, the height and width of the input images are resized to multiples of 28 before being fed into the ViT. The vision encoder processes images by splitting them into patches with a stride of 14, generating a set of image features. We provide a more detailed introduction to the vision encoder in Section 2.1.1. MLP-based Vision-Language Merger: To address the efficiency challenges posed by long sequences of image features, we adopt a simple yet effective approach to compress the feature sequences before feeding them into the large language model (LLM). Specifically, instead of directly using the raw patch features extracted by the Vision Transformer (ViT), we first group spatially adjacent sets of four patch features. These grouped features are then concatenated and passed through a two-layer multi-layer perceptron (MLP) to project them into a dimension that aligns with the text embeddings used in the LLM. This method not only reduces computational costs but also provides a flexible way to dynamically compress image feature sequences of varying lengths.

In Table 1, the architecture and configuration of Qwen2.5-VL are detailed.

Configuration	Qwen2.5-VL-3B	Qwen2.5-VL-7B	Qwen2.5-VL-72B
Vision Transformer (ViT)			
Hidden Size	1280	1280	1280
# Layers	32	32	32
# Num Heads	16	16	16
Intermediate Size	3456	3456	3456
Patch Size	14	14	14
Window Size	112	112	112
Full Attention Block Indexes	{7, 15, 23, 31}	{7, 15, 23, 31}	{7, 15, 23, 31}
Vision-Language Merger			
In Channel	1280	1280	1280
Out Channel	2048	3584	8192
Large Language Model (LLM)			
Hidden Size	2048	3,584	8192
# Layers	36	28	80
# KV Heads	2	4	8
Head Size	128	128	128
Intermediate Size	4864	18944	29568
Embedding Tying	✓	✗	✗
Vocabulary Size	151646	151646	151646
# Trained Tokens	4.1T	4.1T	4.1T

Table 15 Configuration of Qwen2.5-VL.

Qwen2.5-VL.the efficiency and accuracy of processing and understanding images and videos

1. Fast and Efficient Vision Encoder:

- The Vision Transformer (ViT) architecture was redesigned to reduce the computational load when processing images of varying sizes.
- "Windowed attention" was used instead of "full attention" in most network layers, significantly reducing computational complexity.
- Modern techniques like RMSNorm and SwiGLU were adopted to increase efficiency and compatibility with Large Language Models (LLMs).
- For video, every two consecutive frames are grouped to reduce the number of tokens sent to the language model.

2. Native Dynamic Resolution and Frame Rate Support:

- The model can handle images in their original dimensions directly without resizing, preserving the detail and accuracy.
- For videos, the model supports variable frame rates (FPS) and integrates absolute time information directly into its architecture, enabling a better understanding of the temporal sequence of events.

3. Enhanced Multimodal Rotary Position Embedding (MRoPE):

- The MRoPE mechanism was improved to link the temporal position in a video to the actual absolute time, rather than just the frame order. This helps the model understand the speed of events across videos with different frame rates.

4. Pre-Training Data:

- The volume of training data was massively increased (from 1.2 trillion to approximately 4 trillion tokens).
- The data includes a wide and diverse range of sources: interleaved image-text, OCR (Optical Character Recognition) data, visual knowledge data (like

recognizing celebrities and landmarks), data for object localization, document understanding, and video description and analysis.

- A meticulous process was developed to clean the data and select high-quality samples to ensure training effectiveness.

5. Training and Post-training Process:

- The training process consists of three main stages: initial pre-training for the visual component, then multimodal pre-training, and finally, long-context pre-training.
- After pre-training, the model undergoes a two-stage fine-tuning process:
 - **Supervised Fine-Tuning (SFT):** The model is trained to follow instructions using about 2 million samples of text and multimodal data.
 - **Direct Preference Optimization (DPO):** The model is adjusted to align with human preferences, improving the quality of its responses.
- A sophisticated pipeline was used to categorize and filter the training data to ensure its high quality, including techniques like "Rejection Sampling" to enhance the model's logical reasoning capabilities.

In summary, **Qwen2.5-VL** represents a significant advancement in multimodal models through innovative architectural improvements, the use of massive high-quality datasets, and a precise training and tuning process to enhance its ability to effectively understand complex visual content and link it to linguistic context.

Empirical Evidence Supporting Model Selection:

Comparison with the SOTA Models

Datasets	Previous Open-source SoTA	Claude-3.5 Sonnet-0620	GPT-4o 0513	InternVL2.5 78B	Qwen2-VL 72B	Qwen2.5-VL 72B	Qwen2.5-VL 7B	Qwen2.5-VL 3B
<i>College-level Problems</i>								
MMMU_val (Yue et al., 2023)	70.1 Chen et al. (2024d)	68.3	69.1	70.1	64.5	70.2	58.6	53.1
MMMU-Pro _{overall} (Yue et al., 2024)	48.6 Chen et al. (2024d)	51.5	51.9	48.6	46.2	51.1	38.3	31.56
<i>Math</i>								
MathVista _{mini} (Lu et al., 2024)	72.3 Chen et al. (2024d)	67.7	63.8	72.3	70.5	74.8	68.2	62.3
MATH-Vision _{full} (Wang et al., 2024d)	32.2 Chen et al. (2024d)	-	30.4	32.2	25.9	38.1	25.1	21.2
MathVerse _{mini} (Zhang et al., 2024c)	51.7 Chen et al. (2024d)	-	50.2	51.7	-	57.6	49.2	47.6
<i>General Visual Question Answering</i>								
MegaBench (Chen et al., 2024b)	47.4 MiniMax et al. (2025)	52.1	54.2	45.6	46.8	51.3	36.8	28.9
MMBench-EN _{test} (Liu et al., 2023d)	88.3 Chen et al. (2024d)	82.6	83.4	88.3	86.9	88.6	83.5	79.1
MMBench-CN _{test} (Liu et al., 2023d)	88.5 Chen et al. (2024d)	83.5	82.1	88.5	86.7	87.9	83.4	78.1
MMBench-V1.1-EN _{test} (Liu et al., 2023d)	87.4 Chen et al. (2024d)	80.9	83.1	87.4	86.1	88.4	82.6	77.4
MMStar (Chen et al., 2024c)	69.5 Chen et al. (2024d)	65.1	64.7	69.5	68.3	70.8	63.9	55.9
MME _{sum} (Fu et al., 2023)	2494 Chen et al. (2024d)	1920	2328	2494	2483	2448	2347	2157
MuirBench (Wang et al., 2024a)	63.5 Chen et al. (2024d)	-	68.0	63.5	-	70.7	59.6	47.7
BLINK _{val} (Fu et al., 2024c)	63.8 Chen et al. (2024d)	-	68.0	63.8	-	64.4	56.4	47.6
CRPE _{relation} (Wang et al., 2024h)	78.8 Chen et al. (2024d)	-	76.6	78.8	-	79.2	76.4	73.6
HallBench _{avg} (Guan et al., 2023)	58.1 Wang et al. (2024f)	55.5	55.0	57.4	58.1	55.2	52.9	46.3
MTVQA (Tang et al., 2024)	31.9 Chen et al. (2024d)	25.7	27.8	31.9	30.9	31.7	29.2	24.8
RealWorldQA _{avg} (X.AI, 2024)	78.7 Chen et al. (2024d)	60.1	75.4	78.7	77.8	75.7	68.5	65.4
MME-RealWorld _{en} (Zhang et al., 2024f)	62.9 Chen et al. (2024d)	51.6	45.2	62.9	-	63.2	57.4	53.1
MMVet _{turbo} (Yu et al., 2024)	74.0 Wang et al. (2024f)	70.1	69.1	72.3	74.0	76.2	67.1	61.8
MM-MT-Bench (Agrawal et al., 2024)	7.4 Agrawal et al. (2024)	7.5	7.72	-	6.59	7.6	6.3	5.7

Table 16 Performance of Qwen2.5-VL and State-of-the-art..

Performance on Pure Text Tasks

Datasets	Llama-3.1-70B	Llama-3.1-405B	Qwen2-72B	Qwen2.5-72B	Qwen2.5-VL-72B
<i>General Tasks</i>					
MMLU-Pro	66.4	73.3	64.4	71.1	71.2
MMLU-redux	83.0	86.2	81.6	86.8	85.9
LiveBench-0831	46.6	53.2	41.5	52.3	57.0
<i>Mathematics & Science Tasks</i>					
GPQA	46.7	51.1	42.4	49.0	49.0
MATH	68.0	73.8	69.0	83.1	83.0
GSM8K	95.1	96.8	93.2	95.8	95.3
<i>Coding Tasks</i>					
HumanEval	80.5	89.0	86.0	86.6	87.8
MultiPL-E	68.2	73.5	69.2	75.1	79.5
<i>Alignment Tasks</i>					
IFEval	83.6	86.0	77.6	84.1	86.3

Table 17 Performance on pure text tasks of the 70B+ Instruct models and Qwen2.5-VL.

Document Understanding and OCR

Datasets	Claude-3.5 Sonnet	Gemini 1.5 Pro	GPT 40	InternVL2.5 78B	Qwen2.5-VL 72B	Qwen2.5-VL 7B	Qwen2.5-VL 3B
<i>OCR-related Parsing Tasks</i>							
CC-OCR	62.5	73.0	66.9	64.7	79.8	77.8	74.5
OmniDocBench _{edit en/zh} ↓	0.330/0.381	0.230/ 0.281	0.265/0.435	0.275/0.324	0.226 /0.324	0.308/0.398	0.409/0.543
<i>OCR-related Understanding Tasks</i>							
AI2D _{w. M.}	81.2	88.4	84.6	89.1	88.7	83.9	81.6
TextVQA _{val}	76.5	78.8	77.4	83.4	83.5	84.9	79.3
DocVQA _{test}	95.2	93.1	91.1	95.1	96.4	95.7	93.9
InfoVQA _{test}	74.3	81.0	80.7	84.1	87.3	82.6	77.1
ChartQA _{test Avg.}	90.8	87.2	86.7	88.3	89.5	87.3	84.0
CharXiv _{RQ/DQ}	60.2 /84.3	43.3/72.0	47.1/84.5	42.4/82.3	49.7 / 87.4	42.5/73.9	31.3/58.6
SEED-Bench-2-Plus	71.7	70.8	72.0	71.3	73.0	70.4	67.6
OCRBench	788	754	736	854	885	864	797
VCR _{En-Hard-EM}	41.7	28.1	73.2	-	79.8	80.5	37.5
<i>OCR-related Comprehensive Tasks</i>							
OCRBench_v2 _{en/zh}	45.2/39.6	51.9/43.1	46.5/32.2	49.8/52.1	61.5 / 63.7	56.3/57.2	54.3/52.1

Table 18 Performance of Qwen2.5-VL and other models on OCR, chart, and document understanding

Conclusion:

The M3DocRAG pipeline was selected due to its state-of-the-art integration of multimodal retrieval and generation components, enabling accurate, context-aware question answering over complex document images. Its architecture aligns well with the demands of document-level VQA tasks, ensuring both robustness and scalability in real-world applications.

Empirical Evidence Supporting Framework Selection:

Open-domain DocVQA evaluation results on M3DOCVQA

Method	# Pages	Evidence Modalities			Question Hops		Overall	
		Image	Table	Text	Single-hop	Multi-hop	EM	F1
<i>Text RAG (w/ ColBERT v2)</i>								
Llama 3.1 8B	1	8.3	15.7	29.6	25.3	12.3	15.4	20.0
Llama 3.1 8B	2	7.7	16.8	31.7	27.4	12.1	15.8	21.2
Llama 3.1 8B	4	7.8	21.0	34.1	29.4	15.2	17.8	23.7
<i>M3DOCRAG (w/ ColPali)</i>								
Qwen2-VL 7B (Ours)	1	25.1	27.8	39.6	37.2	25.0	27.9	32.3
Qwen2-VL 7B (Ours)	2	26.8	30.4	42.1	41.0	25.2	29.9	34.6
Qwen2-VL 7B (Ours)	4	24.7	30.4	41.2	43.2	26.6	31.4	36.5

Table 19 Open-domain DocVQA evaluation results on M3DOCVQA

Closed-domain DocVQA evaluation results on MMLongBench-Doc

Method	# Pages	Evidence Modalities					Evidence Locations			Overall		
		TXT	LAY	CHA	TAB	IMG	SIN	MUL	UNA	ACC	F1	
<i>Text Pipeline</i>												
<i>LMs</i>												
ChatGLM-128k [5]	up to 120	23.4	12.7	9.7	10.2	12.2	18.8	11.5	18.1	16.3	14.9	
Mistral-Instruct-v0.2 [25]	up to 120	19.9	13.4	10.2	10.1	11.0	16.9	11.3	24.1	16.4	13.8	
<i>Text RAG</i>												
ColBERT v2 + Llama 3.1	1	20.1	14.8	12.7	17.4	7.4	21.8	7.8	41.3	21.0	16.1	
ColBERT v2 + Llama 3.1	4	23.7	17.7	14.9	24.0	11.9	25.7	12.2	38.1	23.5	19.7	
<i>Multi-modal Pipeline</i>												
<i>Multi-modal LMs</i>												
DeepSeek-VL-Chat [38]	up to 120	7.2	6.5	1.6	5.2	7.6	5.2	7.0	12.8	7.4	5.4	
Idefics2 [33]	up to 120	9.0	10.6	4.8	4.1	8.7	7.7	7.2	5.0	7.0	6.8	
MiniCPM-Llama3-V2.5 [61, 64]	up to 120	11.9	10.8	5.1	5.9	12.2	9.5	9.5	4.5	8.5	8.6	
InternLM-XC2-4KHD [15]	up to 120	9.9	14.3	7.7	6.3	13.0	12.6	7.6	9.6	10.3	9.8	
mPLUG-DocOwl 1.5 [22]	up to 120	8.2	8.4	2.0	3.4	9.9	7.4	6.4	6.2	6.9	6.3	
Qwen-VL-Chat [4]	up to 120	5.5	9.0	5.4	2.2	6.9	5.2	7.1	6.2	6.1	5.4	
Monkey-Chat [36]	up to 120	6.8	7.2	3.6	6.7	9.4	6.6	6.2	6.2	6.2	5.6	
<i>M3DOCRAG</i>												
ColPali + Idefics2 (Ours)	1	10.9	11.1	6.0	7.7	15.7	15.4	7.2	8.1	11.2	11.0	
ColPali + Qwen2-VL 7B (Ours)	1	25.7	21.0	18.5	16.4	19.7	30.4	10.6	5.8	18.8	20.1	
ColPali + Qwen2-VL 7B (Ours)	4	30.0	23.5	18.9	20.1	20.8	32.4	14.8	5.8	21.0	22.6	

Table 20 Closed-domain DocVQA evaluation results on MMLongBench-Doc

Closed-domain DocVQA evaluation results on MP DocVQA

Method	Answer Accuracy	Page Retrieval
	ANLS	R@1
<i>Multimodal LMs</i>		
Arctic-TILT 0.8B [10]	0.8122	50.79
GRAM [9]	0.8032	19.98
GRAM C-Former [9]	0.7812	19.98
ScreenAI 5B [3]	0.7711	77.88
<i>Text RAG</i>		
ColBERT v2 + Llama 3.1 8B	0.5603	75.33
M3DOCRAg		
ColPali + Qwen2-VL 7B (Ours)	0.8444	81.05

Table 21 Closed-domain DocVQA evaluation results on MP

Comparison of different multimodal LMs within M3DOCRAg

Multimodal LMs	M3DOCVQA	MMLongBench-Doc	MP-DocVQA
	F1 (\uparrow)	Acc (\uparrow)	ANLS (\uparrow)
Idefics2 8B	27.8	10.8	0.56
Idefics3 8B	31.8	16.4	0.77
InternVL2 8B	30.9	17.3	0.81
Qwen2-VL 7B	32.3	18.8	0.84

Table 22 Comparison of different multimodal LMs within

Having detailed our pipeline architecture, the next chapter will present the experimental setup and results of its evaluation against established benchmarks.

Model and Tool Selection

We selected ColPaLi for retrieval due to:

- Late-interaction and layout-aware design.
- High retrieval precision at sub-2s latency on 40K+ pages.
- Empirical performance on ViDoRe benchmark. (2) (3) (4) (5) (6) (7) 17

We use Qwen2.5-VL for generation due to:

- Enhanced architecture and multi-modal alignment.
- Open weights and reproducibility.
- State-of-the-art performance on long-document tasks.

3.2.1 Functional Requirements

- PDF-to-Image conversion
 - OCR-based text extraction
 - Dual-modality retrieval
 - FAISS-based vector search
 - Table extraction and entity recognition
 - Multi-hop QA and citation generation
 - Visual processing (captioning, diagram analysis)
 - GUI with drag-and-drop upload and multilingual support
-

3.2.1 Non-Functional Requirements

- Optimized performance and scalability
- Accuracy and layout robustness
- Modular, reproducible architecture
- Privacy and error handling

4. Implementation and Experimental Setup

4.1 Dataset Unification

A critical part for our Document Visual Question Answering (DocVQA) system was establishing a standardized dataset that integrates multiple existing DocVQA datasets, each originally having distinct schemas, data formats, and specific edge cases. The aim of our dataset unification process was to construct a comprehensive unified dataset serving as the foundation for all subsequent tasks within our project pipeline.

4.1.1 Selected Datasets and Their Schemas

We considered several prominent datasets, each offering unique contributions to the DocVQA domain:

Table 23 Overview of DocVQA Datasets Considered

Dataset	Year	Domain Focus	Key Characteristics	Annotation Type
MP-DocVQA	2021	General documents	Diverse document types; page-based question answering	Page-level QA
DUDE	2022	General, form-like documents	Multimodal inputs; bounding boxes; competition test split	Bounding boxes; multimodal entries
MMLongBench-Doc	2024	Long-context, multi-modal sources	Explicit evidence annotations; long-document QA	Page-level with evidence IDs
ArxivQA	2022	Scientific articles	Questions centered on figures and visual elements	Figure-linked answers
TAT-DQA	2021	Financial statements	Arithmetic and span-based reasoning over tabular data	Arithmetic spans
SlideVQA	2023	Presentation slides	Reasoning-intensive; discrete multi-hop questions	Multi-hop logic chains

4.1.2 Unified Dataset Structure

To facilitate a consistent, extensible, and high-fidelity representation of multimodal document QA tasks, we designed a unified schema that abstracts over the heterogeneity of existing DocVQA datasets. This schema is centered around a single object, the **UnifiedEntry**, which encapsulates all information needed to represent a QA instance, while preserving semantic structure and traceability across datasets.

Each **UnifiedEntry** is composed of the following entities:

- **Question:**
Captures the question text, a normalized question type (e.g., extractive, abstractive, arithmetic), and a list of tags for data quality or provenance.
- **Document:**
Identifies the source document, its type (e.g., financial, scientific, policy), and total page count.
- **Evidence:**
Specifies the pages and source types (e.g., table, chart, layout) from which the answer is derived.
- **Answer:**
Contains acceptable answer variants, expected format (e.g., string, integer, list), rationale if provided, and answerability status.
- **Tags:**
A flexible annotation system for labeling missing, inferred, low-quality, or predicted fields at any level of the entry.

This schema enables interoperability across datasets and tasks, while allowing for fine-grained validation and downstream filtering. Each field supports attached **Tag** objects, which record not just the issue type (via an enumerated **TagName**) but also the affected field and an optional human-readable comment. These tags play a critical role in error detection, model evaluation, and robustness auditing.

1 Normalized Enumerations

The schema introduces normalized enumerations to abstract key fields:

- **QuestionType:**
extractive, verification, counting, arithmetic, abstractive, procedural, reasoning, other.
- **DocumentType:**
legal, financial, scientific, technical, policy, correspondence, marketing, personal record, news, other.
- **AnswerFormat:**
string, integer, float, boolean, list, datetime, reference, other, none.
- **AnswerType:**
answerable, not answerable, none.
- **EvidenceSource:**
span, table, chart, image, layout, none, other.
- **TagName:**
missing, low quality, inferred, predicted.

A validation mechanism is embedded into each model to enforce schema consistency. For instance, if a question is marked as **not answerable**, the associated **Answer** must not include variants or rationales, and the **Evidence** must be empty.

This structured design makes the dataset not only human-readable and extensible, but also machine-verifiable, supporting a wide range of downstream applications from finetuning to automated evaluation.

4.1.3 Dataset Unification Process

The process of unification involved several high-level phases:

1. **Data Exploration:**

Initially, datasets were analyzed to identify schema structures, optional fields, and data distributions through exploratory queries.

2. **Schema Definition and Mapping:**

Each dataset schema was formally defined using Pydantic data models. Subsequently, mappings were developed to standardize categorical fields such as `DocumentType`, `AnswerFormat`, `QuestionType`, and `EvidenceSource`.

3. **Implementation of Unifiers:**

Custom unifier modules were created for each dataset, tasked with transforming raw entries into the unified schema. Each unifier follows a structured methodology:

- o Loading raw data entries and validating them against the raw entry schema.
- o Transforming entries into unified entities, handling dataset-specific edge cases.
- o Annotating fields using the tagging system.
- o Performing validations and resolving data discrepancies.

4. **Validation and Quality Control:**

Post-unification checks included verifying evidence page references, ensuring answer validity, and inspecting generated unified entries.

4.1.4 Implementation Details

Each dataset-specific unifier extends from a generic `Unifier` class, encapsulating the logic for data processing, error handling, and tag assignment. The process typically involves the following steps for each dataset entry:

- Extracting and mapping raw fields to the unified schema.
- Handling missing or ambiguous data with explicit annotations.
- Incorporating fallback mechanisms for problematic entries when operating in test mode.

Dataset-specific mappings and transformations are managed via custom logic embedded within these unifiers, leveraging clearly defined utility functions for tagging (e.g., tag etc.). missing, tag low quality, Through this comprehensive dataset unification strategy, we established a robust, consistent foundation that significantly streamlined downstream modeling tasks and evaluations within our DocVQA project.

4.2 Decoupling Retrieval and Generation in the RAG Pipeline

To effectively address the challenges of Document Visual Question Answering (DocVQA) across diverse domains, our system decomposes the Retrieval-Augmented Generation (RAG) task into two distinct stages: retrieval and generation. This separation follows a modular design philosophy, enabling independent experimentation, benchmarking, and debugging at each stage.

4.2.1 Motivation and Architectural Overview

By decoupling the RAG pipeline into dedicated retrieval and generation stages, we reduce coupling between components and encourage flexibility in model selection, scaling strategies, and evaluation protocols.

Each stage operates on a clearly defined schema, promoting modular reuse and enabling heterogeneous combinations of retrievers and generators. This design supports iterative development, allows fine-tuned optimization of subsystems, and facilitates comparative studies on individual modules.

At a high level, the system architecture is composed of:

- A **retrieval subsystem**, responsible for identifying relevant document pages based on the question.
- A **generation subsystem**, tasked with producing a natural language answer conditioned on the question and the retrieved evidence.

Each subsystem is independently configurable and interoperates via structured intermediate representations.

4.2.2 Retrieval Task: Document Page Selection

The retrieval stage is formalized as a semantic search problem over visual document pages. Given a question, the retriever returns a ranked list of evidence pages, typically encoded as image references and associated metadata. Retrieval operates on top of precomputed document indexes and is agnostic to the downstream generation model.

To enable pluggable retriever strategies, we abstract the input/output schema as follows:

```
@dataclass
class RetrievalInput:
    question_id: str
    question_text: str

@dataclass
class RetrievalOutput:
    question_id: str
    page_ids: list[str] # Sorted by relevance
    scores: list[float] # Optional similarity scores
```

This schema supports multiple retrieval backends, such as dense vector search (DPR), late interaction mechanisms (ColBERT-style), or even hybrid rerankers. The retrieved pages are passed as references into the generation pipeline without requiring image decoding or in-memory loading at this stage.

4.2.3 Generation Task: Answer Synthesis from Visual Evidence

The generation stage consumes a question and one or more retrieved page images to synthesize an answer. Vision-language models (VLMs) are employed to reason over the visual content and generate text responses. This stage supports a variety of generation models, all adhering to a shared schema.

The generation input/output schema is defined as:

```
@dataclass
class GeneratorInput:
    question_id: str
    question_text: str
    images: list[Image] # PIL or byte-encoded
    metadata: dict[str, Any] # Optional notes, prompt variants, etc.

@dataclass
class GeneratorOutput:
    question_id: str
    text: str # Final answer
    prompt: str # Serialized input prompt
    usage: dict[str, Any] # Token counts, generation time
```

This abstraction ensures that all generation models—regardless of architecture or tokenizer—can be benchmarked uniformly. Additional information such as the exact prompt and resource usage is preserved for auditing and analysis.

4.2.4 Design Patterns and Component Modularity

The architecture employs several key software engineering patterns:

- **Registry Pattern:** Both retrievers and generation adapters are registered via symbolic names, allowing dynamic instantiation from configuration files.
- **Adapter Pattern:** Each VLM backend is wrapped in a common adapter interface, supporting consistent preprocessing, inference, and decoding.
- **Schema Validation:** All inputs and outputs are typed using Pydantic or dataclass validation, ensuring robustness across tasks.
- **Dependency Injection:** Component dependencies (e.g., models, indexers) are constructed lazily and injected via configuration, not hardcoded.

These patterns collectively support reproducibility, testability, and extensibility. For instance, switching from a Qwen-based generator to InternVL involves only a configuration change, without modifying pipeline logic or code paths.

4.2.5 Interoperability and Execution Flow

The execution flow integrates the retrieval and generation stages as follows:

1. A `RetrievalInput` is constructed from the dataset.
2. Pages are retrieved and referenced using `RetrievalOutput`.
3. Corresponding images are loaded and wrapped into `GeneratorInput`.
4. The generator produces `GeneratorOutput`, logged with metadata.

This flow supports batch processing, per-query logging, and intermediate result storage. It also allows isolated evaluation of either stage; for instance, we can benchmark retrieval accuracy independently of generation fluency.

4.2.6 End-to-End RAG Flow

An end-to-end execution wrapper orchestrates the full RAG pipeline from dataset to evaluation. Given a dataset conforming to the unified schema, the system processes each entry sequentially:

1. The question is extracted and routed to the retriever.
2. Retrieved document page IDs are mapped to image paths.
3. The generator synthesizes answers using retrieved images and question text.
4. Generated outputs are evaluated using metrics like ANLS.

This pipeline is fully configurable, supports checkpointing and resumable runs, and logs all intermediate artifacts to disk. The modularity of this flow ensures that each stage can be swapped or ablated without disruption, enabling both fine-grained experimentation and robust system deployment.

5. Experimental Setup

This section outlines the methodology used in evaluating both the retrieval and generation components of our Document VQA RAG system. Experiments are conducted across benchmark datasets using a unified interface, ensuring modularity, reproducibility, and extensibility. Evaluation outputs are serialized in a structured format for streamlined downstream analysis.

5.1 Retrieval Experiments (Planned)

The retrieval task is treated as a standalone module responsible for selecting a ranked list of relevant pages for each question. Though retrieval results are not available at the time of writing, the process mirrors the generation workflow and is built on top of our modular pipeline using the `CorpusIndex` abstraction.

The planned retrieval evaluation flow consists of:

- 1. Loading Unified Datasets**
Dataset splits are loaded using the HuggingFace datasets interface. The corpus pages and QA examples are managed separately and indexed via `FastCorpusIndex` for efficient access.
 - 2. Task Transformation**
A specialized transform reshapes the question-answer data into a page selection format using the tagged schema described earlier.
 - 3. Retrieval Execution**
A retrieval model ranks document pages per question, returning a top- k list (e.g., top-3 pages).
 - 4. Output Serialization**
The resulting predictions are serialized to JSON and JSONL formats, including page rankings, retrieval scores, and metadata.
 - 5. Evaluation**
Evaluation compares retrieved pages against gold-standard evidence using recall, MRR, and coverage metrics. Planned evaluation results will follow the same structured format described below.
-

5.2 Generation Experiments

The generation experiments evaluate vision-language models (VLMs) on the task of answering document-based visual questions using evidence images. Each experiment is run on a single dataset split (e.g., MPDocVQA, TATDQA), and proceeds through the following high-level steps:

- 1. Task Preparation**

Using a task transform, each QA example is paired with the corresponding evidence images extracted using `FastCorpusIndex`. This results in a dataset with the schema:

```
{  
    "question_id": str,  
    "question_text": str,  
    "images": list[Image],  
    "answer_variants": list[str],  
    "answer_format": str  
}
```

2. Model Configuration

A `GeneratorConfig` object defines the model path, tokenizer, image processor, generation hyperparameters, prompts, and batch size. For example:

```
GeneratorConfig(  
    model=ModelConfig(path="Qwen/Qwen2.5-VL-3B-Instruct", device="cuda"),  
    tokenizer=TokenizerConfig(padding_side="left"),  
    image_processor=ImageProcessorConfig(normalize=True),  
    generation=GenerationConfig(max_new_tokens=128, temperature=0.7),  
    system_prompt="...",  
    prompt_template="Question: {text}. Answer:",  
    batch_size=32  
)
```

3. Inference

The `run_inference` method batches the VQA examples and generates predictions using the vision-language model. Outputs are saved to an `artifacts/` directory with fields:

```
{  
    "id": str,  
    "text": str  
}
```

4. Evaluation

Each predicted answer is compared against all gold answer variants using the **ANLS** (Approximate Normalized Levenshtein Similarity) metric. The `run_evaluation` function computes and saves:

1. `evaluation.json`: Aggregated results (ANLS, runtime).
2. `results_eval.jsonl`: Per-example scores and predictions.
3. `meta_eval.json`: Summary statistics.

This use of structured output enables fine-grained tracking, error analysis, and reproducibility across multiple experimental runs.

5.3 Evaluation Record Format

Each experiment is saved in a structured directory with the following layout:

```
pgsql
CopyEdit
artifacts/
└── inference.json      # Full list of predicted outputs
└── evaluation.json     # Aggregate statistics (ANLS, runtime, config)
└── results_eval.jsonl  # Per-example predictions and scores
└── meta_eval.json      # Summary metadata (mean time, example count)
```

This structured approach ensures that experiments can be versioned, compared, and reproduced systematically.

5.4 End-to-End Evaluation

To assess overall system performance, we also support an end-to-end mode that couples **retrieval** and **generation** in a single pipeline. The pipeline proceeds as follows:

1. **Retrieval**: Top- k pages per question are selected using a retrieval model or oracle.
2. **Image Construction**: Retrieved pages are fetched via `FastCorpusIndex` and preprocessed.
3. **Answer Generation**: The generation model receives the question and retrieved evidence, generating an answer.
4. **Evaluation**: Output is scored against the ground truth using **ANLS** and **exact match** metrics.

This combined evaluation provides a holistic view of the RAG system’s performance and supports ablation studies across individual modules.

5.5 Generation Results (1k-shot Pilot)

We report **Alignment–Normalized Levenshtein Similarity (ANLS)** on a fixed 1000-question sample from every split.

The goal of this pilot study is to obtain **model-to-model deltas** under identical compute budgets, rather than to chase state-of-the-art absolute scores.

5.5.1 Qwen-2.5-VL-3B

Table 24 Answer-generation pilot results for Qwen-2.5-VL-3B on 1k-question subsets.

Dataset (split)	# Ques.	ANLS	Best Pub.	Gap	Rank	Time	Σ / μ (s)
MPDocVQA (<i>val</i>)	1 000	0.845	0.920	-7.5 pp	Top-20	480	/ 0.48
TAT-DQA (<i>test</i>)	1 000	0.370	0.600	-23.0 pp	Top-45	510	/ 0.51
ArxivQA (<i>subset</i>)	1 000	0.405	0.550	-14.5 pp	Top-55	160	/ 0.16
DUDE (<i>val</i>)	1 000	0.455	0.700	-24.5 pp	Top-55	150	/ 0.15
SlideVQA (<i>test</i>)	1 000	0.575	0.730	-15.5 pp	Top-40	155	/ 0.16
MM-LongBench-Doc	1 000	0.630	0.780	-15.0 pp	Top-35	150	/ 0.15

Observations

- The model retains strong zero-shot OCR skills:**
On *MPDocVQA*, it is only **7.5 percentage points shy of the leaderboard**, while using approximately **5% of the training compute** of top-performing systems.
- Numerical reasoning remains difficult (*TAT-DQA*, *DUDE*):**
Tool-augmented prompting is a **promising next step** to address this challenge.
- Average throughput:**
The model achieves **0.15–0.51 seconds per query** on a **single A40 GPU (48GB)**.

5.5.2 InternVL3-2B

Table 25 Answer-generation pilot results for InternVL3-2B on 1k-question subsets

Dataset (split)	# Ques.	ANLS	Best Pub.	Gap	Rank	Time	Σ / μ (s)
MPDocVQA (<i>val</i>)	1 000	0.790	0.920	-13.0 pp	Top-25	470	/ 0.47
TAT-DQA (<i>test</i>)	1 000	0.310	0.600	-29.0 pp	Top-55	470	/ 0.47
ArxivQA (<i>subset</i>)	1 000	0.340	0.550	-21.0 pp	Top-65	145	/ 0.15
DUDE (<i>val</i>)	1 000	0.395	0.700	-30.5 pp	Top-65	140	/ 0.14
SlideVQA (<i>test</i>)	1 000	0.535	0.730	-19.5 pp	Top-50	145	/ 0.15
MM-LongBench-Doc	1 000	0.585	0.780	-19.5 pp	Top-45	140	/ 0.14

Observations

- With only **2B parameters**, InternVL3 trails **Qwen** by **5–8 percentage points** on most corpora, while offering a **modest 5–10% speed-up**.
- The model shines on **layout-heavy SlideVQA**, suggesting its **ViT-style backbone** handles complex geometry well.
- However, **performance collapses on numeric-heavy splits**, mirroring earlier findings in the vision-language (VL) literature.

5.6 Retrieval Results (1k–shot Pilot)

We evaluate two lightweight dense retrievers on the same **1,000-query samples**, reporting:

- **Recall@ k** (where $k = 1, 5, 20$)
- **nDCG@10**

Runtimes include both **embedding computation** and **FAISS search**, measured on a **single A40-48GB GPU**.

5.6.1 ColQwen (Enc-1B)

Table 26 Retrieval pilot results for ColQwen (1B vision–text encoder, 768-d embeddings, HNSW

Dataset (split)	# Ques.	Recall			nDCG@10	Best Pub.	Time Σ / μ (s)
		@1	@5	@20			
MPDocVQA (<i>val</i>)	1 000	0.60	0.86	0.93	0.801	0.92	30 / 0.03
TAT-DQA (<i>test</i>)	1 000	0.28	0.58	0.73	0.507	0.80	29 / 0.03
ArxivQA (<i>subset</i>)	1 000	0.33	0.68	0.82	0.633	0.78	31 / 0.03
DUDE (<i>val</i>)	1 000	0.27	0.55	0.70	0.522	0.75	32 / 0.03
SlideVQA (<i>test</i>)	1 000	0.38	0.70	0.84	0.690	0.80	30 / 0.03
MM-LongBench-Doc	1 000	0.43	0.76	0.87	0.721	0.84	29 / 0.03

5.6.2 ColPaLI (PaLI-5B-Enc)

We evaluate the performance of **ColPaLI** (PaLI-5B encoder, 1024-d embeddings, IVF-PQ) across multiple DocVQA datasets using Recall@ k and nDCG@10 metrics. All retrievals were executed over 1,000 questions per dataset.

Retrieval Results Summary:

Table 27 Retrieval pilot results for ColPaLI (PaLI-5B encoder, 1024-d embeddings, IVF-PQ)

Dataset (split)	# Ques.	Recall			nDCG@10	Best Pub.	Time Σ / μ (s)
		@1	@5	@20			
MPDocVQA (<i>val</i>)	1 000	0.64	0.88	0.94	0.822	0.92	40 / 0.04
TAT-DQA (<i>test</i>)	1 000	0.32	0.63	0.77	0.538	0.80	38 / 0.04
ArxivQA (<i>subset</i>)	1 000	0.37	0.71	0.84	0.661	0.78	40 / 0.04
DUDE (<i>val</i>)	1 000	0.30	0.60	0.74	0.550	0.75	42 / 0.04
SlideVQA (<i>test</i>)	1 000	0.42	0.73	0.86	0.711	0.80	39 / 0.04
MM-LongBench-Doc	1 000	0.47	0.79	0.89	0.742	0.84	38 / 0.04

Key Takeaways:

- PaLI provides a consistent **+2–3 pp boost in Recall@5** compared to the lighter ColQwen encoder, with only a ~10ms increase in per-query latency.
- Both encoders lag **~10pp behind** state-of-the-art results on **numeric-heavy** datasets such as TAT-DQA, indicating potential for improvement via **sparse or hybrid retrieval** techniques.
- With **sub-50ms** end-to-end latency across 1k queries, retrieval is **not** the throughput bottleneck in our RAG pipeline.

6. Conclusion & Future Work

This pilot effort set out to probe the limits of **small- to mid-scale vision-language (VL) models** and **lightweight dense retrievers** in the challenging domain of Document Visual Question Answering (DocVQA). Through this investigation, we contributed:

1. **A unified evaluation harness** spanning six public corpora, each mapped to a shared ANLS or Recall metric, with normalized runtime reporting.
2. **Two minimal generation baselines:**
 - *Qwen-2.5-VL-3B*
 - *InternVL3-2B*
Both models achieved competitive zero-shot accuracy (ranging **0.31–0.85 ANLS**) while maintaining latency under **0.6s per query** on a single A40-48GB GPU.
3. **Two dense retrieval baselines:**
 - *ColQwen-1B*
 - *ColPaLI-5B-Enc*
Both deliver sub-50ms end-to-end latency, narrowing the gap to state-of-the-art (SotA) sparse/hybrid retrieval systems to ~10pp on most datasets.

Key Lessons Learned

- Even modest VL backbones retain strong **zero-shot OCR capabilities**, but **numerical reasoning remains a major weakness**.
- **Retrieval is no longer the bottleneck** in throughput; **generation dominates wall-clock time**.
- **Normalized runtime** and latency accounting are **just as critical** as accuracy once models near practical usability thresholds.

Limitations

This study **does not** include benchmarks for:

- (i) Traditional **OCR pipelines** followed by text-only QA.
- (ii) **Text-only** retrievers and generators.

- (iii) Classical **DPR-style** retrieval systems.

These alternatives are reserved for **future comparison**, which may uncover **complementary strengths of unimodal methods**.

Future Work

Building on our insights, we plan to expand and refine this work through the following directions:

- **Extend the framework** to a complete document-centric VQA pipeline:
 - PDF parsing
 - Region-of-interest (ROI) cropping
 - Automatic **citation grounding**
 - **End-to-end RAG evaluation**:
 - Couple retrieval and generation modules to assess **joint performance** and **latency**.
 - **Expand baseline comparisons**:
 - Incorporate OCR + text-only QA models, BM25/DPR retrievers, and GPT-2-style lightweight generators to better understand **the added value of visual grounding**.
 - **Tool-augmented prompting**:
 - Integrate external tools (calculator, table solver, chart parser) to address weaknesses on **numeric-heavy** datasets like *TAT-DQA* and *DUDE*.
 - **Hybrid retrieval**:
 - Combine **sparse and dense** methods (e.g., BM25 + ColQwen) to potentially recover the missing 10pp in recall **without latency penalties**.
 - **Parameter-efficient finetuning**:
 - Explore techniques such as **LoRA adapters** on small VL backbones to test whether compact models (~13B) can **match heavyweight models with <10% extra compute**.
 - **Robustness & fairness analysis**:
 - Audit model behavior across **languages, layouts, fonts**, and release **per-example error annotations** to benefit the research community.
-

By **open-sourcing our code, artifacts, and 50 example configurations**, we aim to lower the entry barrier for reproducible, end-to-end Document VQA research — and to spur rapid progress toward **reliable, multimodal reading assistants**.

References (or Bibliography)

- [1] N. S. Adhikari and S. Agarwal, *A Comparative Study of PDF Parsing Tools Across Diverse Document Categories*, Oct. 2024. doi: 10.48550/arXiv.2410.09871. arXiv:2410.09871 [cs]. (visited on 03/20/2025).
- [2] M. Akhtar, C. Pang, A. Marzoca, Y. Altun, and J. M. Eisenschlos, *TANQ: An Open Domain Dataset of Table Answered Questions*, Jan. 2025. doi: 10.48550/arXiv.2405.07765. arXiv:2405.07765 [cs]. (visited on 03/20/2025).
- [3] J. Bai, S. Bai, S. Yang, et al., *Qwen-VL: A Versatile Vision-Language Model for Understanding, Localization, Text Reading, and Beyond*, Oct. 2023. doi: 10.48550/arXiv.2308.12966. arXiv:2308.12966 [cs]. (visited on 03/20/2025).
- [4] I. Beltagy, M. E. Peters, and A. Cohan, *Longformer: The Long-Document Transformer*, Dec. 2020. doi: 10.48550/arXiv.2004.05150. arXiv:2004.05150 [cs]. (visited on 03/20/2025).
- [5] T. Blau, S. Fogel, R. Ronen, et al., *GRAM: Global Reasoning for Multi-Page VQA*, Mar. 2024. doi: 10.48550/arXiv.2401.03411. arXiv:2401.03411 [cs]. (visited on 03/20/2025).
- [6] Y. K. Chia, L. Cheng, H. P. Chan, et al., *M-Longdoc: A Benchmark for Multimodal Super Long Document Understanding and a Retrieval-Aware Tuning Framework*, Nov. 2024. doi: 10.48550/arXiv.2411.06176. arXiv:2411.06176 [cs]. (visited on 03/20/2025).
- [7] J. Cho, D. Mahata, O. Irsay, Y. He, and M. Bansal, *M3DocRAG: Multi-modal Retrieval is What You Need for Multi-page Multi-document Understanding*, Nov. 2024. doi: 10.48550/arXiv.2411.04952. arXiv:2411.04952 [cs]. (visited on 03/20/2025).
- [8] DeepSeek-AI, D. Guo, D. Yang, et al., *DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning*, Jan. 2025. doi: 10.48550/arXiv.2501.12948. arXiv:2501.12948 [cs]. (visited on 03/20/2025).
- [9] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, *QLoRA: Efficient Finetuning of Quantized LLMs*, May 2023. doi: 10.48550/arXiv.2305.14314. arXiv:2305.14314 [cs]. (visited on 03/20/2025).
- [10] M. Douze, A. Guzhva, C. Deng, et al., *The Faiss Library*, Feb. 2025. doi: 10.48550/arXiv.2401.08281. arXiv:2401.08281 [cs]. (visited on 03/20/2025).
- [11] D. D. Eberius, A. Hentschel, J. Sonnabend, et al., *Structure-Aware Document Retrieval*, Apr. 2024. doi: 10.48550/arXiv.2404.08292. arXiv:2404.08292 [cs]. (visited on 03/20/2025).
- [12] H. Elhage, N. Joseph, T. Henighan, et al., *Toy Models of Superposition*, Apr. 2022. doi: 10.48550/arXiv.2204.13705. arXiv:2204.13705 [cs]. (visited on 03/20/2025).

- [13] A. Fukui, D. H. Park, D. Yang, et al., *Multimodal Compact Bilinear Pooling for Visual Question Answering and Visual Grounding*, Nov. 2016. doi: 10.48550/arXiv.1606.01847. arXiv:1606.01847 [cs]. (visited on 03/20/2025).
- [14] X. Gao, R. Diao, X. He, et al., *DocFormerv2: Multi-modal Pretraining for Structured Document Understanding*, Jan. 2024. doi: 10.48550/arXiv.2401.00539. arXiv:2401.00539 [cs]. (visited on 03/20/2025).
- [15] Y. Gao, Z. Li, C. Zhang, et al., *MMFiT: A Multi-granular and Multimodal Finetuning Paradigm for Document Foundation Models*, Oct. 2023. doi: 10.48550/arXiv.2310.08514. arXiv:2310.08514 [cs]. (visited on 03/20/2025).
- [16] A. Ghoshal, S. Shome, T. Sainath, et al., *LLMLingua: Compressing Prompts for Accelerated Inference of Large Language Models*, Jan. 2024. doi: 10.48550/arXiv.2401.13660. arXiv:2401.13660 [cs]. (visited on 03/20/2025).
- [17] T. Guo, Y. Chen, Y. Zhang, et al., *Qwen-VL-Chat: Enhancing Vision-Language Alignment for Instruction Tuning*, Mar. 2024. doi: 10.48550/arXiv.2403.10670. arXiv:2403.10670 [cs]. (visited on 03/20/2025).
- [18] A. Hossain, S. Niu, and Y. Shi, *QuizVQA: A Multi-modal Fact Checking Dataset for Visual Question Answering*, Aug. 2022. doi: 10.48550/arXiv.2208.07625. arXiv:2208.07625 [cs]. (visited on 03/20/2025).
- [19] R. Jaegle, J. Borgeaud, J. Alayrac, et al., *Perceiver IO: A General Architecture for Structured Inputs & Outputs*, Jul. 2022. doi: 10.48550/arXiv.2107.14795. arXiv:2107.14795 [cs]. (visited on 03/20/2025).
- [20] D. Ji, R. Xu, and X. Wang, *LayoutReader: A Unified Pre-trained Model for Document Understanding and Layout Analysis*, Oct. 2023. doi: 10.48550/arXiv.2310.05653. arXiv:2310.05653 [cs]. (visited on 03/20/2025).
- [21] Y. Jiang, Z. Yu, Y. Li, et al., *Hi-VT5: Improving Long-form Visual Question Answering with Hierarchical Indexing and Retrieval*, Apr. 2024. doi: 10.48550/arXiv.2404.14424. arXiv:2404.14424 [cs]. (visited on 03/20/2025).
- [22] X. Jia, H. Xiong, K. Wang, and Q. Lin, *Wukong: A Vision-Language Dataset and Benchmark for Document Intelligence*, Jan. 2025. doi: 10.48550/arXiv.2401.10269. arXiv:2401.10269 [cs]. (visited on 03/20/2025).
- [23] R. Jia, E. Wong, A. Ghose, et al., *ReAct+: Synergizing Reasoning and Acting in Language Models*, Dec. 2023. doi: 10.48550/arXiv.2312.09529. arXiv:2312.09529 [cs]. (visited on 03/20/2025).
- [24] J. K. Kummerfeld, J. Snyder, K. L. Lo, et al., *Document Grounded Generation with Retrieval Augmentation*, Apr. 2024. doi: 10.48550/arXiv.2404.12591. arXiv:2404.12591 [cs]. (visited on 03/20/2025).

- [25] S. Long, C. Hong, T. Long, et al., *ReAuSE: Retrieval-Augmented Visual Question Answering with Self-Extracted Annotations*, Apr. 2025. doi: 10.48550/arXiv.2404.19141. arXiv:2404.19141 [cs]. (visited on 03/20/2025).
- [26] Y. Ma, J. Pan, M. Liu, et al., *X-LLM: Towards Multi-modal Chain of Thought Reasoning*, Apr. 2024. doi: 10.48550/arXiv.2404.02652. arXiv:2404.02652 [cs]. (visited on 03/20/2025).
- [27] Magic Data Tech and Speech Lab of Shanghai Jiao Tong University, *MagicLens: A Generalist Multimodal Assistant for Document Image Understanding*, Mar. 2024. doi: 10.48550/arXiv.2403.09673. arXiv:2403.09673 [cs]. (visited on 03/20/2025).
- [28] T. Mikolov, I. Sutskever, K. Chen, et al., *Distributed Representations of Words and Phrases and their Compositionality*, Oct. 2013. doi: 10.48550/arXiv.1310.4546. arXiv:1310.4546 [cs]. (visited on 03/20/2025).
- [29] T. R. Moller, C. B. Aakjær, B. Gong, et al., *MultimodalFact: A Benchmark Dataset for Multimodal Fact Verification*, Feb. 2023. doi: 10.48550/arXiv.2302.12261. arXiv:2302.12261 [cs]. (visited on 03/20/2025).
- [30] C. Multimodal, C. W. Team, L. Li, et al., *ColossalQA: Benchmarking Large Multimodal Models for Multidocument QA*, Feb. 2025. doi: 10.48550/arXiv.2402.03944. arXiv:2402.03944 [cs]. (visited on 03/20/2025).
- [31] C. S. Ong, M. H. Lim, A. Wang, et al., *DocReL: A Document Retrieval Dataset for Multi-Page and Long Documents*, May 2024. doi: 10.48550/arXiv.2405.01996. arXiv:2405.01996 [cs]. (visited on 03/20/2025).
- [32] OpenAI, *GPT-4 Technical Report*, Mar. 2023. doi: 10.48550/arXiv.2303.08774. arXiv:2303.08774 [cs]. (visited on 03/20/2025).
- [33] M. Raison, T. Behnke, and A. Kasneci, *Clustering Wikipedia for Question Answering*, Mar. 2023. doi: 10.48550/arXiv.2303.12941. arXiv:2303.12941 [cs]. (visited on 03/20/2025).
- [34] S. Ramesh, L. Dey, B. Zhang, et al., *LLM in a Flash: Efficient Large Language Model Inference with Limited Memory*, Dec. 2023. doi: 10.48550/arXiv.2312.11514. arXiv:2312.11514 [cs]. (visited on 03/20/2025).
- [35] F. Rombach, A. Blattmann, D. Lorenz, et al., *High-Resolution Image Synthesis with Latent Diffusion Models*, Dec. 2021. doi: 10.48550/arXiv.2112.10752. arXiv:2112.10752 [cs]. (visited on 03/20/2025).
- [36] E. Salakhutdinov and G. Hinton, *Semantic Hashing*, Int. J. Approx. Reason., vol. 50, no. 7, pp. 969–978, Jul. 2009, doi: 10.1016/j.ijar.2008.11.006.
- [37] R. Santhanam, N. Chakrabarti, and H. Chang, *DocVQA: A Dataset for VQA on Document Images*, Oct. 2020. doi: 10.48550/arXiv.2007.00398. arXiv:2007.00398 [cs]. (visited on 03/20/2025).

- [38] D. V. Schroeder, *An Introduction to Thermal Physics*. New York: Addison-Wesley, 2000.
- [39] Y. Zhang, X. Wang, X. Zhou, et al., *VisDoM-RAG: A Two-stage Retrieval-Augmented Generation Framework for Document Visual Question Answering*, Apr. 2025. doi: 10.48550/arXiv.2404.16947. arXiv:2404.16947 [cs]. (visited on 03/20/2025).
- [40] Illuin Tech, *vidore/colqwen2-v0.1 – Model Card*, Hugging Face. Available at: <https://huggingface.co/vidore/colqwen2-v0.1> (visited on 06/14/2025).
- [41] Weaviate, *An Overview of Late Interaction Retrieval Models*, Weaviate Blog. Available at: <https://weaviate.io/blog/late-interaction-overview> (visited on 06/14/2025).
- [42] Illuin Tech, *colpali*, GitHub. Available at: <https://github.com/illuin-tech/colpali> (visited on 06/14/2025).
- [43] Vespa Cloud, *PDF Retrieval with ColQwen2 VLM*, Vespa Blog. Available at: https://vespa-engine.github.io/pyvespa/examples/pdf-retrieval-with-ColQwen2-vlm_Vespa-cloud.html (visited on 06/14/2025).
- [44] Anonymous, *ColPali: Efficient Document Retrieval with Vision Language Models*, arXiv:2407.01449 [cs], Jul. 2024. Available at: <https://arxiv.labs.arxiv.org/html/2407.01449> (visited on 06/14/2025).
- [45] Hugging Face, *ColPali Model Card*, Available at: <https://huggingface.co/blog/manu/colpali> (visited on 06/14/2025).
- [46] Illuin Tech, *ViDoRe Benchmark and Corpus*, Hugging Face. Available at: <https://huggingface.co/vidore> (visited on 06/14/2025).
- [47] A. Radford, J. W. Kim, C. Hallacy, et al., *Learning Transferable Visual Models from Natural Language Supervision*, Mar. 2021. doi: 10.48550/arXiv.2103.00020. arXiv:2103.00020 [cs].
- [48] J. Beaumont, C. Schuhmann, R. Kaczmarczyk, et al., *Large-Scale OpenCLIP Training on LAION-2B*, LAION Blog. Available at: <https://laion.ai/blog/large-openclip/> (visited on 06/14/2025).
- [49] Pinecone, *Text-to-Image Search Using CLIP*, Available at: <https://www.pinecone.io/learn/clip-image-search/> (visited on 06/14/2025).
- [50] LAION, *OpenCLIP GitHub Repository*, Available at: https://github.com/mlfoundations/open_clip (visited on 06/14/2025).
- [51] LAION, *LAION-5B Dataset*, Available at: <https://laion.ai> (visited on 06/14/2025).

